


Part 1:

Solving Problems with Search

[Acknowledgment: Some Slides adapted from Dan Klein and Pieter Abbeel]


<http://ai.berkeley.edu>]

Search




Today

- Agents that Plan Ahead
- Search Problems
- Search Algorithms (Review?)
 - Depth-First Search
 - Breadth-First Search
 - Iterative Deepening



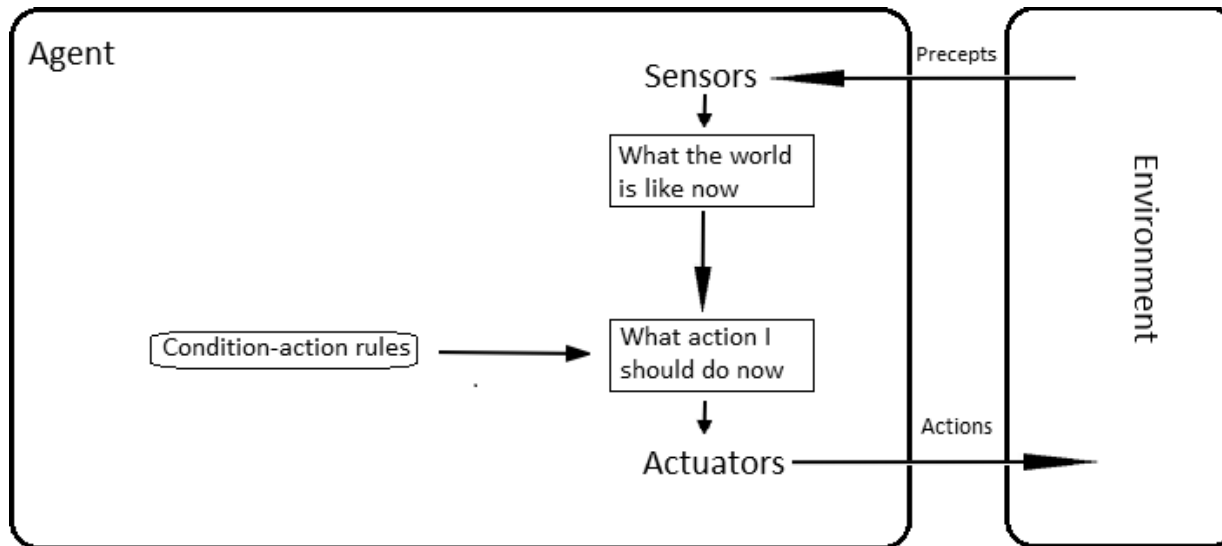
Agents that Plan



Agent types


- Five basic types in order of increasing generality:
 - Table Driven agents
 - **Simple reflex agents (you will make one now)**
 - Model-based reflex agents
 - Goal-based agents
 - **Problem-solving agents**
 - Utility-based agents
 - **Can distinguish between different goals**
 - Learning agents

Reflex Agents



- Act only on the basis of the current percept.
- Based on the *condition-action rules*:
 if condition
 then action
- **Infinite loops are common** (can be fixed with some randomization)

Example: Vacuum Agent (R&N, 2.1)



Percept sequence	Action
[A, Clean]	Right
[A, Dirty]	Suck
[B, Clean]	Left
[B, Dirty]	Suck
[A, Clean], [A, Clean]	Right
[A, Clean], [A, Dirty]	Suck
⋮	⋮
[A, Clean], [A, Clean], [A, Clean]	Right
[A, Clean], [A, Clean], [A, Dirty]	Suck
⋮	⋮

Vacuum Reflex agent (Python)

```
def ReflexVacuumAgent():  
    "A reflex agent for the two-state vacuum environment."  
    def program(percept):  
        location, status = percept  
        if status == 'Dirty':  
            return 'Suck'  
        elif location == loc_A:  
            return 'Right'  
        elif location == loc_B:  
            return 'Left'  
    return Agent(program)
```

<https://github.com/aimacode/aima-python/blob/master/agents.py>

Reflex Agent for Pacman

```
python pacman.py --layout testMaze --pacman GoWestAgent
```

But, things get ugly for this agent when turning is required:

```
python pacman.py --layout tinyMaze --pacman GoWestAgent
```



In-Class Exercise: WallFollower Agent

- Step 1: download search.zip from Canvas
- Step 2: find searchAgents.py
- Step 3: Copy code below, rename to new class


```
class GoWestAgent(Agent): → WallFollower
    "An agent that goes west until it can't."
    def getAction(self, state):
        "The agent receives a GameState (defined in pacman.py)."
        if Directions.WEST in state.getLegalPacmanActions():
            return Directions.WEST
        else:
            return Directions.STOP
```

Reflex Agents in General

- Reflex agents:
 - Choose action based on current percept (and maybe memory)
 - May have memory or a model of the world's current state
 - Do not consider the future consequences of their actions
 - Consider how the world IS
- Can a reflex agent be rational?




from



<http://ai.berkeley.edu/>

Planning Agents

- Planning agents:
 - Ask “what if”
 - Decisions based on (hypothesized) consequences of actions
 - Must have a model of how the world evolves in response to actions
 - Must formulate a goal (test)
 - **Consider how the world
WOULD BE**
- Optimal vs. complete planning
- Planning vs. replanning



Rational agents

- **Performance measure:** An objective criterion for success of an agent's behavior, e.g.,
 - Robot driver?
 - Chess-playing program?
 - Spam email classifier?
- **Rational Agent:** selects actions that is *expected* to **maximize its performance measure**,
 - given percept sequence
 - given agent's built-in knowledge
 - consider: how to maximize expected future performance, with only historical data?

Task Environment

- Before we design an intelligent agent, we must specify its “task environment”:

PEAS:

Performance measure

Environment

Actuators

Sensors

PEAS: Example

- Example: Agent = robot driver in DARPA Challenge

- Performance measure:

-

- Environment:

-

- Actuators:

-

- Sensors:

-



Stanley

PEAS: Example

- Example: Agent = robot driver in DARPA Challenge
 - Performance measure:
 - Time to complete course
 - Environment:
 - Road, obstacles
 - Actuators:
 - Steering wheel, accelerator, brake, signal, horn
 - Sensors:
 - Optical cameras, lasers, sonar, accelerometer, speedometer, GPS, odometer, engine sensors, ...




Goal-Based Agents: Search

- Five basic types in order of increasing generality:
 - Table Driven agents
 - Simple reflex agents
 - Model-based reflex agents
 - Goal-based agents
 - Problem-solving agents: solve problems by searching for solution
 - Utility-based agents
 - Learning agents


Search Problems

- A **search problem** consists of:

- State space



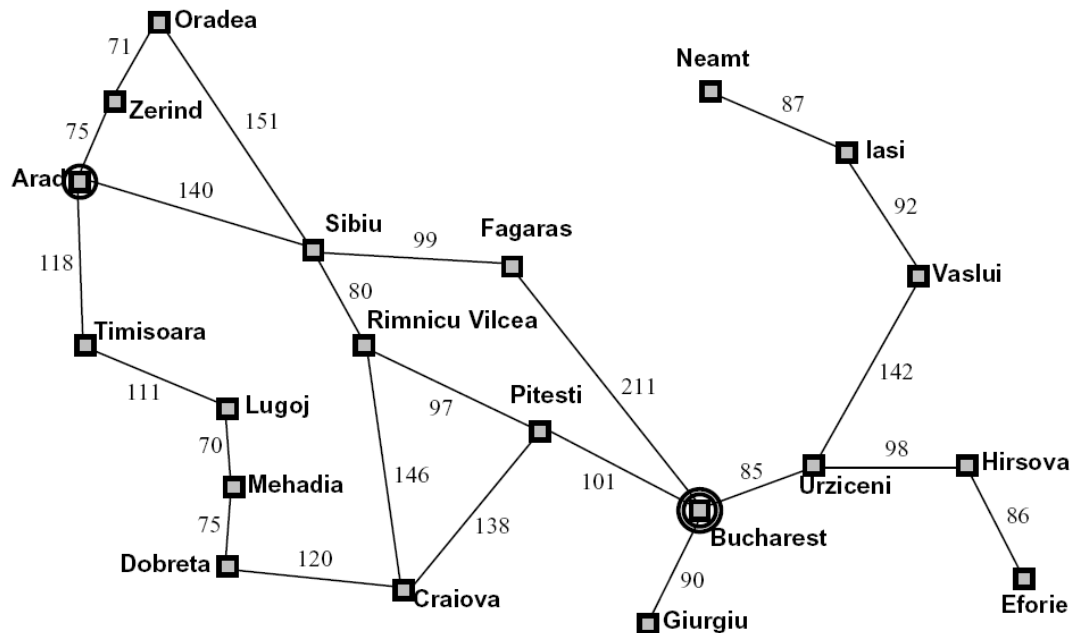
- Successor function
(with actions, costs)



- Start state and a goal test

- A **solution** is a sequence of **actions** which transforms the start state to a goal state

Example: Traveling in Romania



- State space:
—
- Successor function:
—
- Start state:
—
- Goal test:
—
- Solution?

Example: 8-puzzle

- states?
- initial state?
- Successor function?
- goal test?
- solution?

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Path finding: Amazon PrimeAir™


- First successful delivery in Cambridge, U.K in December
- Fill in the blank:
 - states?
 - initial state?
 - Successor function?
 - goal test?
 - solution?



<https://www.youtube.com/watch?v=vNySOrl2Ny8>

How to represent State Space?

The **world state** includes every last detail of the environment




A **search state** keeps **only** the details needed for planning (abstraction).

Note: *Static info (e.g., walls) do not need to be stored for each state*


- Problem: Path-Finding
 - States: (x,y) location
 - Actions: NSEW
 - Successor: update location only
 - Goal test: is $(x,y)=END$
- Problem: Eat-All-Dots
 - States: $\{(x,y), \text{dots: bool array}\}$
 - Actions: NSEW
 - Successor: update location and possibly dots
 - Goal test: dots all false

State Space Sizes?

- World state:
 - Agent positions: 120
 - Food count: 30
 - Ghost positions: 12
 - Agent direction: NSEW
- How many
 - World states?
 $120 \times (2^{30}) \times (12^2) \times 4$
 - States for path-finding?
120
 - States for eat-all-dots?
 $120 \times (2^{30})$




Exercise: Safe Passage



- Problem: eat all dots while keeping the ghosts scared
- What does the state space have to specify?


State Space Graphs: 1

- **State space graph:** A mathematical representation of a search problem
 - Nodes are (abstracted) world configurations
 - Arcs represent successors (action results)
 - The goal test is a set of goal nodes (maybe only one)
- In a state space graph, each state occurs only once!
 - Careful: if there are loops in this graph, keep track of visited states
- We can rarely build this full graph in memory (it's too big), but it's a useful idea




State Space Graphs: 2

- State space graph: A mathematical representation of a search problem
 - Nodes are (abstracted) world configurations
 - Arcs represent successors (action results)
 - The goal test is a set of goal nodes (maybe only one)
- In a search graph, each state occurs only once!
- We can rarely build this full graph in memory (it's too big), but it's a useful idea



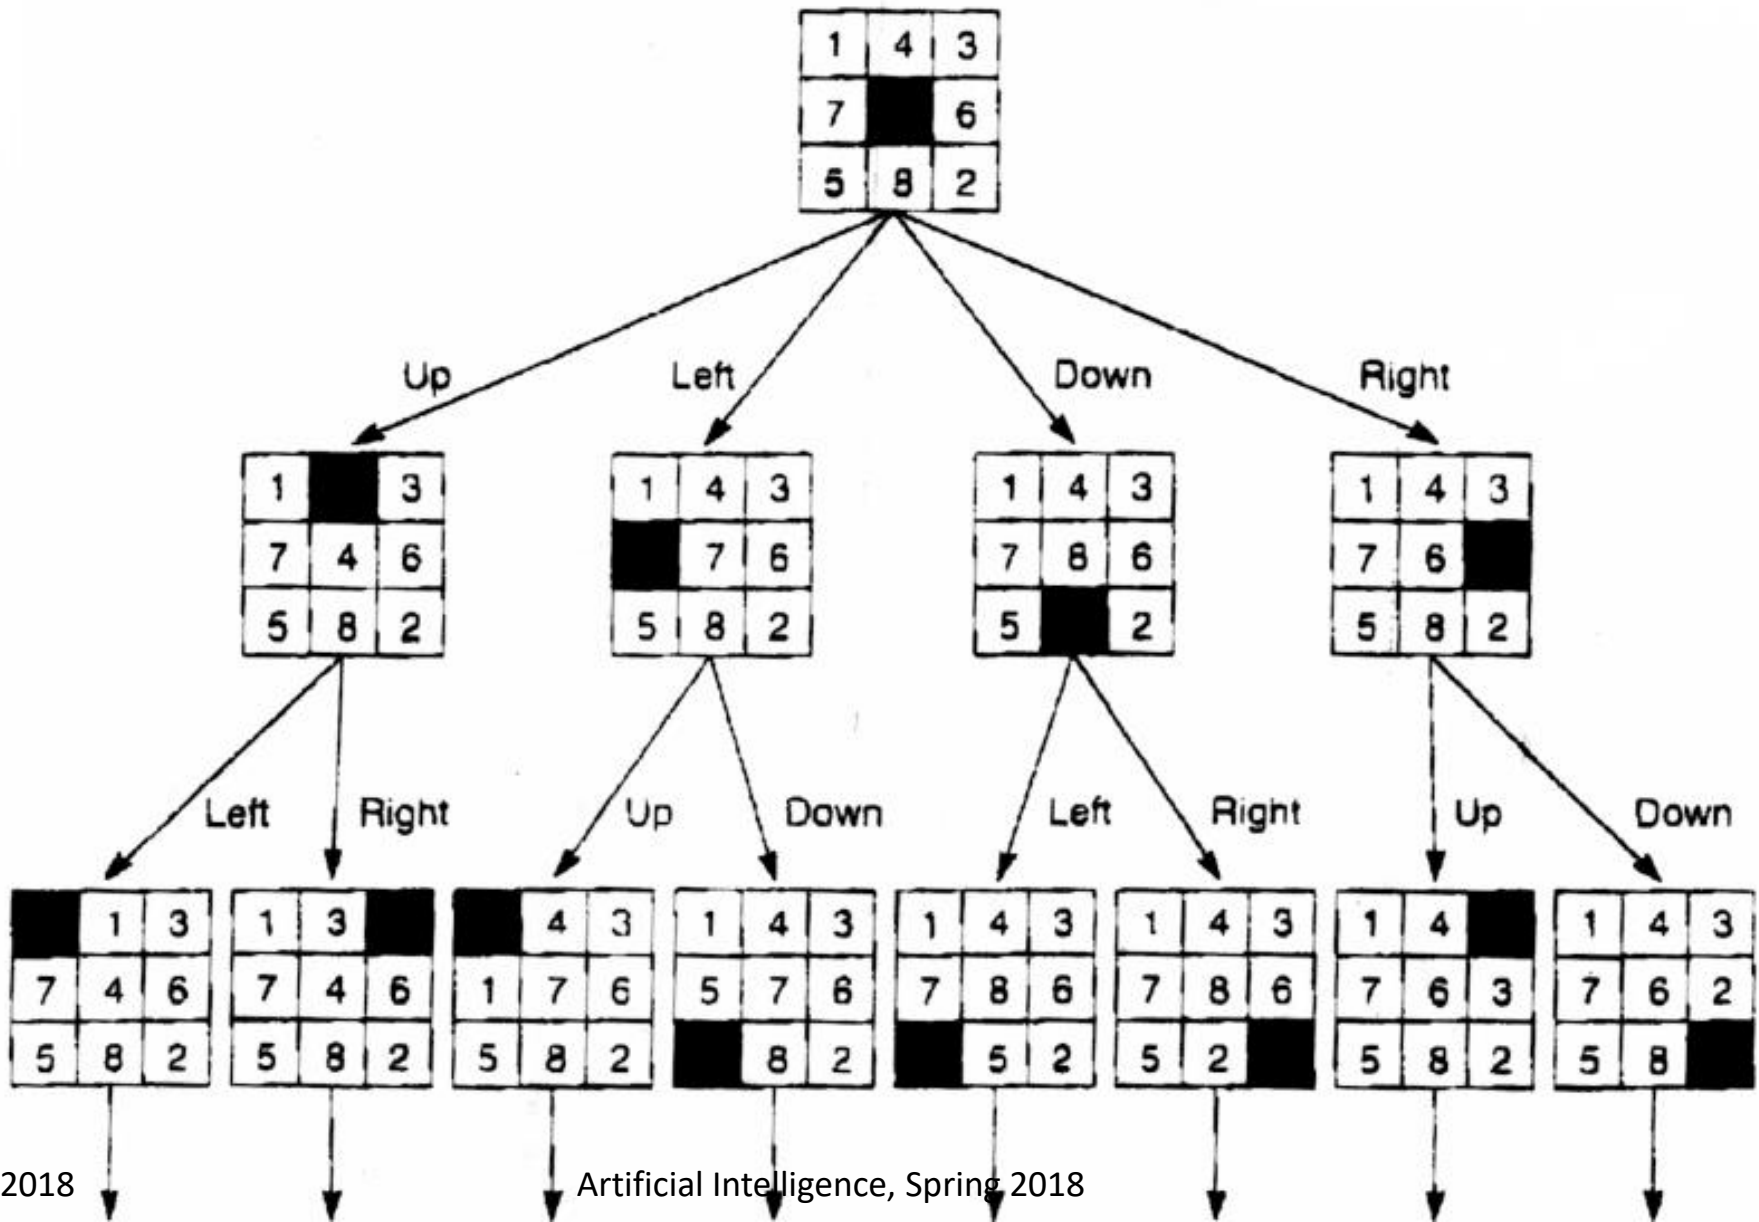
Tiny search graph for a tiny search problem

Search Trees




- A search tree:
 - A “what if” tree of plans and their outcomes
 - The start state is the root node
 - Children correspond to successors
 - Nodes show states, but correspond to **PLANS (paths from root to the state)**
 - For most problems, we can never actually build the whole tree (too large)

Search Tree for 8 puzzle problem



State Space Graphs vs. Search Trees


State Space Graph



Each **NODE** in in the search tree represents **entire PATH** in the state space graph.


We construct both on demand – and we construct as little as possible.

Search Tree




Exercise: State Space Graphs vs. Search Trees

Consider this 4-state graph:




Draw the Search Tree:



Exercise: State Space Graphs vs. Search Trees

Consider this 4-state graph:




How big is the Search Tree?

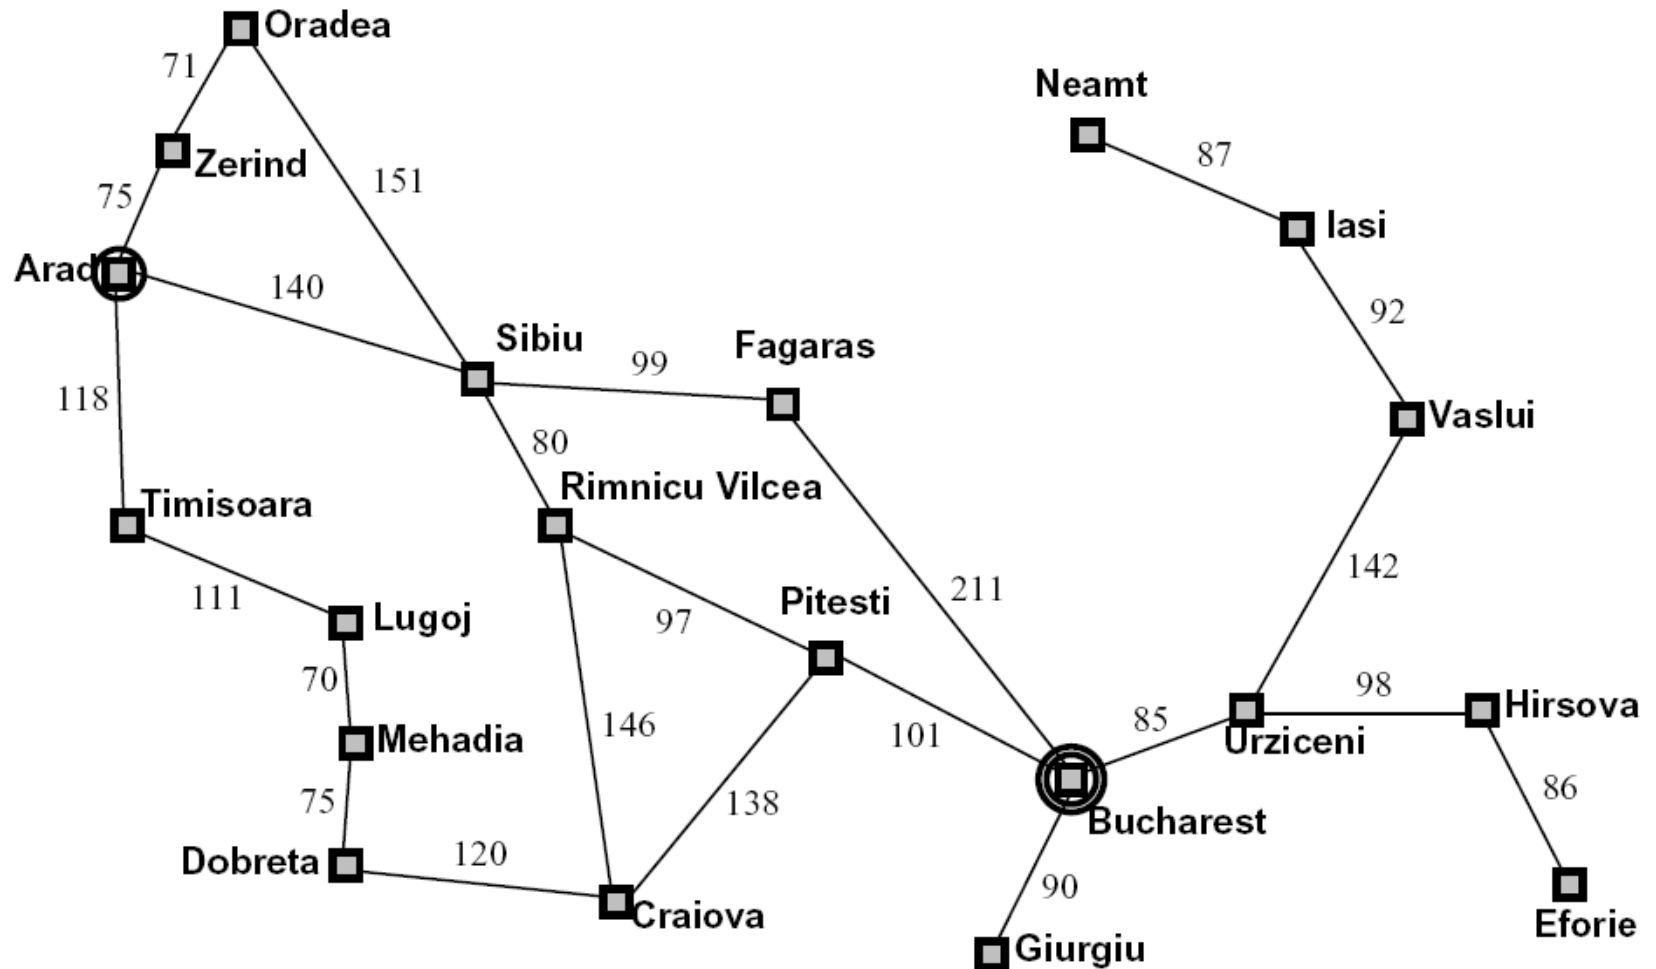


Problem: Lots of repeated structures in the search tree!


Tree Search



Search Example: Romania



Searching with a Search Tree



- Search:
 - Expand out potential plans (tree nodes)
 - Maintain a **fringe** of partial plans under consideration
 - Try to expand as few tree nodes as possible

General Tree Search

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```

- Important ideas:
 - Fringe
 - Expansion
 - Exploration strategy
- Main question: which fringe nodes to explore?

Python Implementation

```
def tree_search(problem, fringe):  
    """Search through the successors of a problem to find a goal.  
    The argument fringe should be an empty queue.  
    Don't worry about repeated paths to a state. [Figure 3.7]"""  
  
    fringe.append( Node(problem.initial) )  
  
    while fringe: //a.k.a: fringe.len()>0  
        node = fringe.pop()  
  
        if problem.goal_test(node.state):  
            return node  
  
        fringe.extend(node.expand(problem) )  
  
    return None
```

adapted from <https://github.com/aimacode/aima-python/blob/master/search.py>

Why Search can be hard

Assuming $b=10$, 1000 nodes/sec, 100 bytes/node

Depth of Solution	Nodes to Expand	Time	Memory
0	1	1 millisecond	100 bytes
2	111	0.1 seconds	11 kbytes
4	11,111	11 seconds	1 megabyte
8	10^8	31 hours	11 giabytes
12	10^{12}	35 years	111 terabytes



$$P(40) \approx \frac{64!}{32! (8!)^2 (2!)^6} \approx 10^{43}.$$

Sidebar: Search vs. Intuition



- Human chess grandmasters think “only” 3-5 moves ahead (Kasparov occasionally 12-14) but rely on patterns, intuition
- Deep blue and others: exhaustive search for optimal state/solution. Evaluates 100M positions/sec, vs. Kasparov 3 positions/sec

Search Strategies: Notation

- A **search strategy** is defined by picking the order of node expansion (fringe exploration)
- Strategies are evaluated along the following dimensions:
 - **completeness**: does it always find a solution if one exists?
 - **time complexity**: number of nodes generated
 - **space complexity**: maximum number of nodes in memory
 - **optimality**: does it always find a least-cost solution?
- Time and space complexity are measured in terms of
 - **b**: maximum branching factor of the search tree
 - **d**: depth of the least-cost solution
 - **m**: maximum depth of the state space (may be ∞)



Depth-First Search



Depth-First Search


*Strategy: expand
a deepest node
first*

*Implementation:
Fringe is a LIFO
stack*




Search Algorithm Properties

- Complete: Guaranteed to find a solution if one exists?
- Optimal: Guaranteed to find the least cost path?
- Time complexity?
- Space complexity?
- Cartoon of search tree:
 - b is the branching factor
 - m is the maximum depth
 - solutions at various depths
- Number of nodes in entire tree?
 - $1 + b + b^2 + \dots + b^m = O(b^{m+1})$



Depth-First Search (DFS) Properties


- What nodes DFS expand?
 - Some left prefix of the tree.
 - Could process the whole tree!
 - If m is finite, takes time $O(b^m)$
- How much space does the fringe take?
 - Only has siblings on path to root, so $O(bm)$
- Is it complete?
 - m could be infinite, so yes iff we prevent cycles (more later)
- Is it optimal?
 - No, it finds the “leftmost” solution, regardless of depth or cost



Project 1: Pacman Search

<http://www.mathcs.emory.edu/~eugene/cs425/p1/>

Due: Friday February 9th



To Do: Start reading Project 1 code

Files you'll edit:

[search.py](#)

Where all of your search algorithms will reside.

[searchAgents.py](#)

Where all of your search-based agents will reside.

Files you might want to look at:

[pacman.py](#)

The main file that runs Pac-Man games. This file describes a Pac-Man GameState type, which you use in this project.

[game.py](#)

The logic behind how the Pac-Man world works. This file describes several supporting types like AgentState, Agent, Direction, and Grid.

[util.py](#)

Useful data structures for implementing search algorithms