


Part 1:

Solving Problems with Search

[Acknowledgment: Some Slides adapted from Dan Klein and Pieter Abbeel]

<http://ai.berkeley.edu>]


Search, continued




Search Problems

- A **search problem** consists of:

- State space



- Successor function
(with actions, costs)




- Start state and a goal test

- A **solution** is a sequence of **actions** which transforms the start state to a goal state


State Space Graphs: 1

- **State space graph:** A mathematical representation of a search problem
 - Nodes are (abstracted) world configurations
 - Arcs represent successors (action results)
 - The goal test is a set of goal nodes (maybe only one)
- In a state space graph, each state occurs only once!
 - Careful: if there are loops in this graph, keep track of visited states
- We can rarely build this full graph in memory (it's too big), but it's a useful idea




State Space Graphs: 2

- State space graph: A mathematical representation of a search problem
 - Nodes are (abstracted) world configurations
 - Arcs represent successors (action results)
 - The goal test is a set of goal nodes (maybe only one)
- In a search graph, each state occurs only once!
- We can rarely build this full graph in memory (it's too big), but it's a useful idea




Tiny search graph for a tiny search problem

Search Trees



- A search tree:
 - A “what if” tree of plans and their outcomes
 - The start state is the root node
 - Children correspond to successors
 - Nodes show states, but correspond to **PLANS (paths from root to the state)**
 - For most problems, we can never actually build the whole tree (too large)

Searching with a Search Tree



- Search:
 - Expand out potential plans (tree nodes)
 - Maintain a **fringe** of partial plans under consideration
 - Try to expand as few tree nodes as possible

General Tree Search

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```

- Important ideas:
 - Fringe
 - Expansion
 - Exploration strategy
- Main question: which fringe nodes to explore?

Python Implementation

```
1 def tree_search(problem, fringe):  
    """Search through the successors of a problem to find a goal.  
    The argument fringe should be an empty queue.  
    Don't worry about repeated paths to a state. [Figure 3.7]"""  
  
2     fringe.append( Node(problem.initial) )  
  
3     while fringe: //a.k.a: fringe.len()>0  
4         node = fringe.pop()  
  
5         if problem.goal_test(node.state):  
6             return node  
  
7         fringe.extend(node.expand(problem) )  
  
9     return None
```

adapted from <https://github.com/aimacode/aima-python/blob/master/search.py>

Search Strategies: Notation

- A **search strategy** is defined by picking the order of node expansion (fringe exploration)
- Strategies are evaluated along the following dimensions:
 - **completeness**: does it always find a solution if one exists?
 - **time complexity**: number of nodes generated
 - **space complexity**: maximum number of nodes in memory
 - **optimality**: does it always find a least-cost solution?
- Time and space complexity are measured in terms of
 - **b**: maximum branching factor of the search tree
 - **d**: depth of the least-cost solution
 - **m**: maximum depth of the state space (may be ∞)



Depth-First Search



Depth-First Search

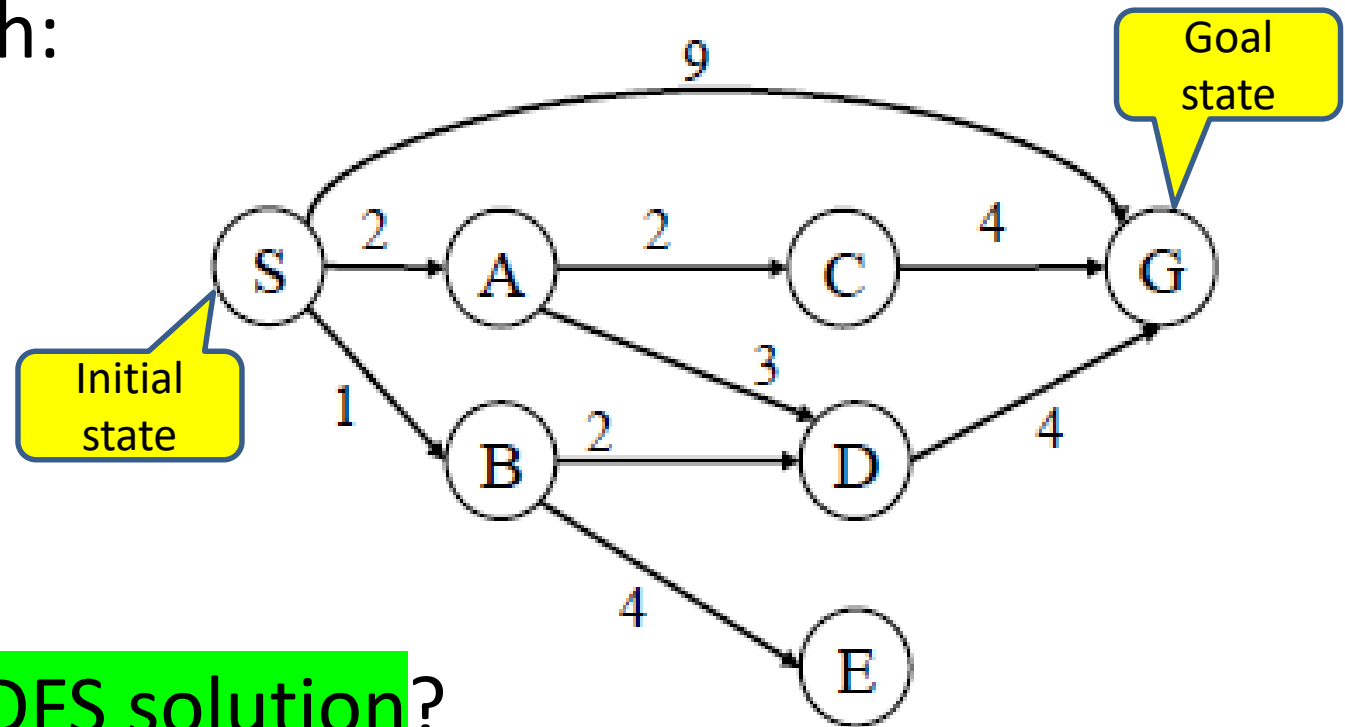
Strategy: expand a
deepest node first

Implementation:
Fringe is a LIFO stack



Exercise: Solve graph by DFS

- Given a graph:



- What is the DFS solution?
 - Assume children stored in alphabetical order.
- Solution: on board

Exercise 2: make into DFS tree search

```
1 def tree_search(problem, fringe):  
    """Search through the successors of a problem to find a goal.  
    The argument fringe should be an empty queue.  
    Don't worry about repeated paths to a state. [Figure 3.7]"""  
  
2     fringe.append( Node(problem.initial) )  
  
3     while fringe: //a.k.a: fringe.len()>0  
4         node = fringe.pop()  
  
5         if problem.goal_test(node.state):  
6             return node  
  
7         fringe.extend(node.expand(problem) )  
  
9     return None
```

adapted from <https://github.com/aimacode/aima-python/blob/master/search.py>


DFS... in 1 line 😊

```
def DFS(problem) :  
    """Search the deepest nodes in the search tree first."""  
    return tree_search(problem, Stack())
```

adapted from <https://github.com/aimacode/aima-python/blob/master/search.py>


Search Algorithm Properties

- Complete: Guaranteed to find a solution if one exists?
- Optimal: Guaranteed to find the least cost path?
- Time complexity?
- Space complexity?
- Cartoon of search tree:
 - b is the branching factor
 - m is the maximum depth
 - solutions at various depths
- Number of nodes in entire tree?
 - $1 + b + b^2 + \dots + b^m = O(b^{m+1})$




Depth-First Search (DFS) Properties

- What nodes DFS expand?
 - Some left prefix of the tree.
 - Could process the whole tree!
 - If m is finite, takes time $O(b^m)$
- How much space does the fringe take?
 - Only has siblings on path to root, so $O(bm)$
- Is it complete?
 - m could be infinite, so yes iff we prevent cycles (more later)
- Is it optimal?
 - No, it finds the “leftmost” solution, regardless of depth or cost





Breadth-First Search



Breadth-First Search

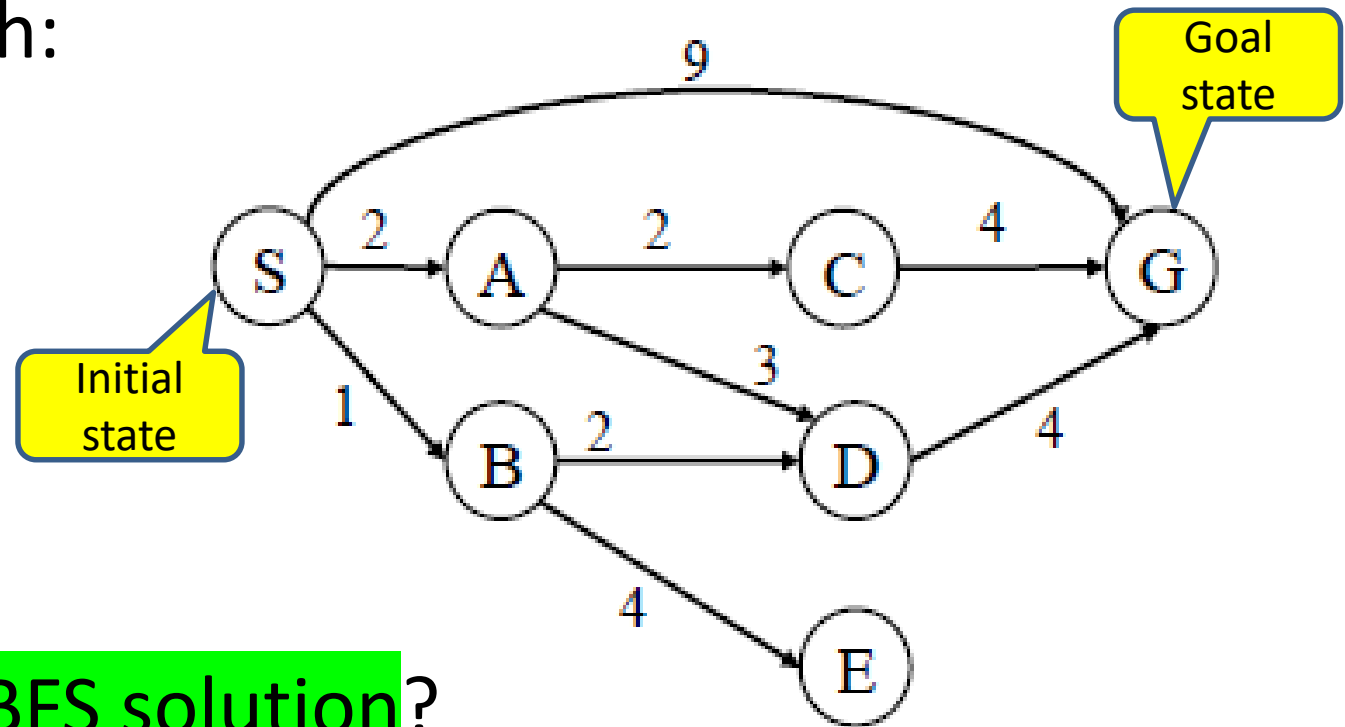
Strategy: expand
a shallowest
node first

Implementation:
Fringe is a FIFO
queue



Exercise: Solve graph by BFS

- Given a graph:



- What is the BFS solution?
 - Assume children stored in alphabetical order.
- Solution: on board

Exercise 3: make into **BFS** tree search

```
1 def tree_search(problem, fringe):  
    """Search through the successors of a problem to find a goal.  
    The argument fringe should be an empty queue.  
    Don't worry about repeated paths to a state. [Figure 3.7]"""  
  
2     fringe.append( Node(problem.initial) )  
  
3     while fringe: //a.k.a: fringe.len()>0  
4         node = fringe.pop()  
  
5         if problem.goal_test(node.state):  
6             return node  
  
7         fringe.extend(node.expand(problem) )  
  
9     return None
```

adapted from <https://github.com/aimacode/aima-python/blob/master/search.py>


BFS ... in 1 line 😊

```
def BFS(problem):  
    """Search the shallowest nodes in the search tree first."""  
    return tree_search(problem, FIFOQueue())
```

adapted from <https://github.com/aimacode/aima-python/blob/master/search.py>

Breadth-First Search (BFS) Properties

- What nodes does BFS expand?
 - Processes all nodes above shallowest solution
 - Let depth of shallowest solution be s
 - Search takes time $O(b^s)$
- How much space does the fringe take?
 - Has roughly the last tier, so $O(b^s)$
- Is it complete?
 - s must be finite if a solution exists, so yes!
- Is it optimal?
 - Only if costs are all 1 (more on costs later)



Memory a Limitation?

- Suppose:
 - 4 GHz CPU
 - 6 GB main memory
 - 100 instructions / expansion
 - 5 bytes / node
- 400,000 expansions / sec
 - Memory filled in 300 sec ... 5 min

Remember: BFS needs to keep $O(b^d)$ states (fringe) in memory

Exercise: DFS vs BFS


- When will BFS outperform DFS?
- When will DFS outperform BFS?

Comparisons


- When will BFS outperform DFS?
- When will DFS outperform BFS?

Algorithm		Complete	Optimal	Time	Space
DFS	w/ Path Checking	Y	N	$O(b^m)$	$O(bm)$
BFS		Y	Y*	$O(b^d)$	$O(b^d)$

Video of Demo Maze Water DFS or BFS? (1)




Video of Demo Maze Water DFS or BFS? (2)



BFS+DFS: Iterative Deepening (ID)

- Idea: combine DFS space advantage with BFS time / shallow-solution advantages
 - Run a DFS with depth limit 1. If no solution...
 - Run a DFS with depth limit 2. If no solution...
 - Run a DFS with depth limit 3.
- Isn't that wastefully redundant?
 - Generally most work happens in the lowest level searched, so not so bad!



Speed

Assuming 10M nodes/sec & sufficient memory

	BFS			Iter. Deep.	
	Nodes	Time		Nodes	Time
8 Puzzle	10^5	.01 sec		10^5	.01 sec
2x2x2 Rubik's	10^6	.2 sec		10^6	.2 sec
15 Puzzle	10^{13}	6 days	1Mx	10^{17}	20k yrs
3x3x3 Rubik's	10^{19}	68k yrs	8x	10^{20}	574k yrs
24 Puzzle	10^{25}	12B yrs		10^{37}	10^{23} yrs

Why the difference?

Rubik has higher branch factor
15 puzzle has greater depth



of duplicates

Slide adapted from Richard Korf presentation




Danger!

- Failure to detect repeated states can cause exponentially more work.



Graph Search: BFS

- In BFS, we shouldn't bother expanding the circled nodes (why?)



Graph Search: **Implementation**

- Idea: never **expand** a state twice
- How to implement:
 - Tree search + set of expanded states (“closed set”)
 - Expand the search tree node-by-node, but...
 - **Before expanding a node**, check to make sure its state has never been expanded before
 - If expanded: skip it, if new: add to closed set
- **Efficiency tip:** **store the closed set as a set**, not a list (why?)

Graph Search PseudoCode

General code.

Exercise: How to make it behave like BFS? DFS?

```
function GRAPH-SEARCH(problem, fringe) return a solution, or failure
  closed ← an empty set
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then return node

    ...

    for child-node in EXPAND(STATE[node], problem) do
      fringe ← INSERT(child-node, fringe)
    end
  end
end
```

Exercise: Python Implementation

```
def graph_search(problem, fringe):  
    """Search through the successors of a problem to find a goal.  
    The argument fringe should be an empty queue.  
    Don't worry about repeated paths to a state. [Figure 3.7]"""  
  
    fringe.append( Node(problem.initial) )  
  
    while fringe:    //a.k.a: fringe.len()>0  
        node = fringe.pop()  
  
        if problem.goal_test(node.state):  
            return node  
  
        fringe.extend(node.expand(problem) )  
  
    return None
```

adapted from <https://github.com/aimacode/aima-python/blob/master/search.py>

Exercise: Python Implementation

```
1 def graph_search(problem, fringe):  
    """Search through the successors of a problem to find a goal.  
    The argument fringe should be an empty queue.  
    closed = {}  
    fringe.append( Node(problem.initial) )  
  
2    while fringe: //a.k.a: fringe.len()>0  
3        node = fringe.pop()  
4  
5        if problem.goal_test(node.state):  
6            return node  
7        if node.state not in closed:  
            closed[node.state]=1  
        else continue  
  
9        fringe.extend(node.expand(problem) )  
  
    return None
```

Node for Search Tree (Python)

```
class Node:
```

```
    """node in a search tree. Contains pointer to parent (the node
    that this is a successor of) and to the actual state for this node.
    Note that if a state is arrived at by two paths, then there are two
    nodes with the same state. Also includes the action that got us to
    this state, and the total path_cost (also known as g) to reach the
    node."""
```

```
    def __init__(self, state, parent=None, action=None, path_cost=0):
        """Create a search tree Node, from a parent by an action."""
        self.state = state
        self.parent = parent
        self.action = action
        self.depth = 0
        if parent:
            self.depth = parent.depth + 1
```

adapted from <https://github.com/aimacode/aima-python/blob/master/search.py>

Node: methods

```
class Node:
    ....
1  def expand(self, problem):
    """List the nodes reachable in one step from this node."""
2      return [ self.successor(problem, action)
3                for action in problem.actions(self.state) ]

4  def successor(self, problem, action):
5      next = problem.result(self.state, action)
6      return Node(next, self, action)

7  def path(self):
8      node, path_back = self, []
9      while node:
10         path_back.append(node)
11         node = node.parent
12     return list( reversed(path_back) )
```

adapted from <https://github.com/aimacode/aima-python/blob/master/search.py>

Node: using in Queue

```
class Node:
```

```
...
```

```
# We want for a queue of nodes in breadth_first_search or  
# DFS to have no duplicated states, so we treat nodes  
# with the same state as equal. [Note: this may not be what you  
# want in other contexts.]
```


```
def __eq__(self, other):
```

```
    return isinstance(other, Node) and self.state == other.state
```

```
def __hash__(self):
```

```
    return hash(self.state)
```


Cost-Sensitive Search (UCS)





BFS finds the shortest path in terms of number of actions.
It does not find the least-cost path. We will now cover
a similar algorithm which does find the least-cost path.


Uniform Cost Search

Strategy: **expand a cheapest node first:**

Fringe is **a priority queue** (priority: cumulative cost)

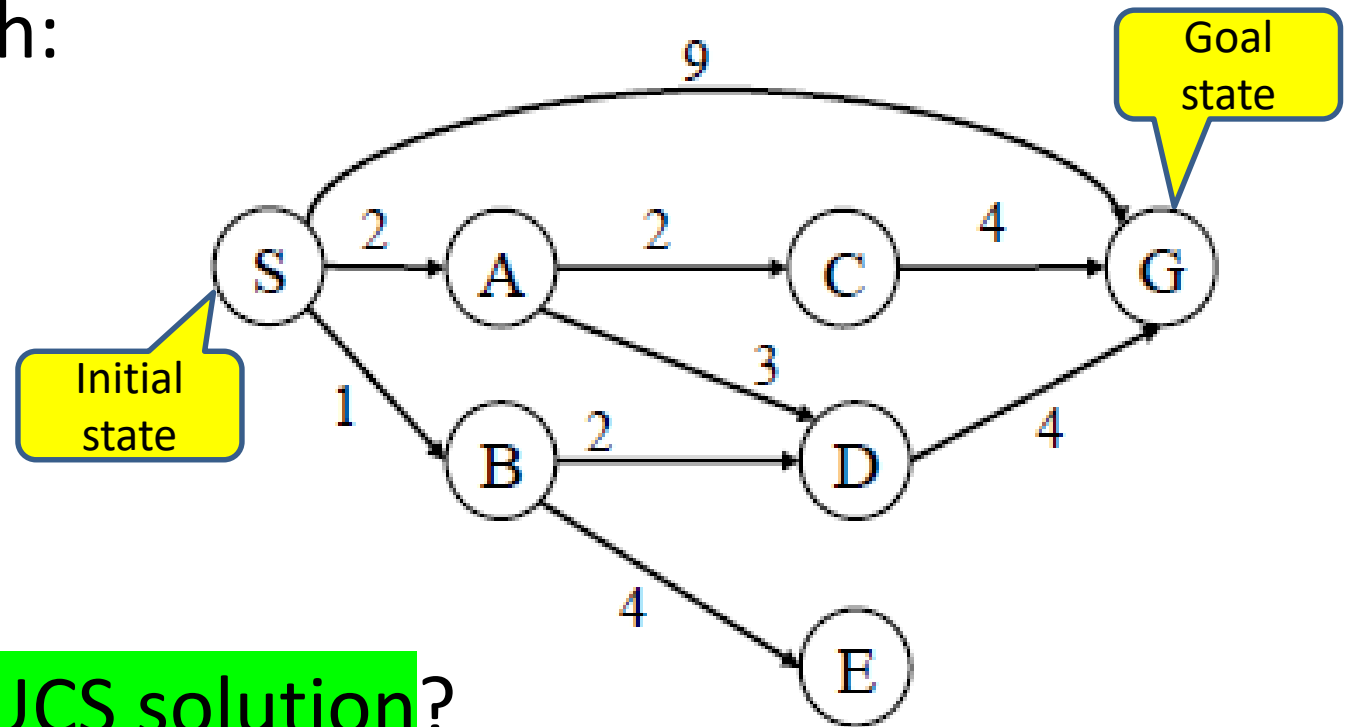


Video of Demo Empty UCS



Exercise: Solve graph by UCS

- Given a graph:



- What is the UCS solution?
 - Assume children stored in alphabetical order.
- Solution: on board

Exercise: make into UCS tree search

```
1 def tree_search(problem, fringe):  
    """Search through the successors of a problem to find a goal.  
    The argument fringe should be an empty queue.  
    Don't worry about repeated paths to a state. [Figure 3.7]"""  
  
2     fringe.append( Node(problem.initial) )  
  
3     while fringe: //a.k.a: fringe.len()>0  
4         node = fringe.pop()  
  
5         if problem.goal_test(node.state):  
6             return node  
  
7         fringe.extend(node.expand(problem) )  
  
9     return None
```

adapted from <https://github.com/aimacode/aima-python/blob/master/search.py>

UCS ... in 1 line 😊

```
def BFS(problem):  
    """Search the shallowest nodes in the search tree first."""  
    return tree_search(problem,  
        util.PriorityQueueWithFunction(Node.getCost))
```



HUH?

<http://www.mathcs.emory.edu/~eugene/cs425/p1/docs/util.html>

Node for UCS (expanded)

```
class Node:
```

```
    def __init__(self, state, parent=None, action=None, path_cost=0):
        """Create a search tree Node, from a parent by an action."""
        self.state = state
        self.parent = parent
        self.action = action
        self.path_cost = path_cost
        self.depth = 0
        if parent:
            self.depth = parent.depth + 1
```


```
    def getCost(self):
        return self.cost
```

adapted from <https://github.com/aimacode/aima-python/blob/master/search.py>



One Queue to rule them all...

- All these search algorithms are the same except for fringe strategies
 - Conceptually, **all fringes are priority queues (i.e. collections of nodes with attached priorities)**
 - Practically, for DFS and BFS, you can avoid the $\log(n)$ overhead from an actual priority queue, by using stack and queues
 - Python Hint: can make one general graph search implementation that takes a variable **Fringe** object as a parameter
 - Use `utils.pm` for Stack, Queue, PriorityQueue classes.





Priority Queue Refresher

- A priority queue is a data structure in which you can insert and retrieve (key, value) pairs with the following operations:


<code>pq.push(key, value)</code>	inserts (key, value) into the queue.
<code>pq.pop()</code>	returns the key with the lowest value, and removes it from the queue.

- You can decrease a key's priority by pushing it again
- Unlike a regular queue, insertions aren't constant time, usually $O(\log n)$
- We'll need priority queues for cost-sensitive search methods

Project 1: Pacman Search

<http://www.mathcs.emory.edu/~eugene/cs425/p1/>

Due: Friday February 9th



Project 1 Tips

- Use Discussions, read FAQ before posting questions:
- **Questions 1-3:** if you develop a correct solution for DFS, the rest will be easy modifications
- **Do not use shortcuts:** Do use Node class or similar. You will need these for Questions 5-8.
- Questions 5-8: more fun/creative. Leave enough time, start early.
- **Most importantly:** Don't Panic! Eat the elephant one question at a time.

To Do: Start solving Q1-3

Files you'll edit:

[search.py](#)

Where all of your search algorithms will reside.

[searchAgents.py](#)

Where all of your search-based agents will reside.

Files you might want to look at:

[pacman.py](#)

The main file that runs Pac-Man games. This file describes a Pac-Man GameState type, which you use in this project.

[game.py](#)

The logic behind how the Pac-Man world works. This file describes several supporting types like AgentState, Agent, Direction, and Grid.

[util.py](#)

Useful data structures for implementing search algorithms