


# Part 1:

# Solving Problems with Search


[Acknowledgment: Some Slides adapted from Dan Klein and Pieter Abbeel]

<http://ai.berkeley.edu>]



Tesla on **autopilot** crashes into **\*parked\*** Fire truck at 65mph



# Search, continued



# Uniform Cost Search (Review)

Strategy: **expand a cheapest node first:**



Fringe is **a priority queue** (priority: cumulative cost)



# UCS Search (Reminder)


Strategy: **expand a cheapest node first:**

Fringe is **a priority queue** (priority: cumulative cost)



## UCS over Maze with Deep/Shallow Water

---



# UCS ... in 1 line 😊

```
def BFS(problem):  
    """Search the shallowest nodes in the search tree first."""  
    return tree_search(problem,  
        util.PriorityQueueWithFunction(Node.getCost))
```




HUH?

<http://www.mathcs.emory.edu/~eugene/cs425/p1/docs/util.html>



# One Queue to rule them all...

- All these search algorithms are the same except for fringe strategies
  - Conceptually, **all fringes are priority queues (i.e. collections of nodes with attached priorities)**
  - Practically, for DFS and BFS, you can avoid the  $\log(n)$  overhead from an actual priority queue, by using stack and queues
  - Python Hint: can make one general graph search implementation that takes a variable **Fringe** object as a parameter
  - Use `utils.pm` for Stack, Queue, PriorityQueue classes.







# Priority Queue Refresher

---


- A priority queue is a data structure in which you can insert and retrieve (key, value) pairs with the following operations:

<code>pq.push(key, value)</code>	inserts (key, value) into the queue.
<code>pq.pop()</code>	returns the key with the lowest value, and removes it from the queue.

- You can decrease a key's priority by pushing it again
- Unlike a regular queue, insertions aren't constant time, usually  $O(\log n)$
- We'll need priority queues for cost-sensitive search methods


# Uniform Cost Search (UCS) Properties

- What nodes does UCS expand?
  - Processes all nodes with cost less than cheapest solution!
  - If that solution costs  $C^*$  and arcs cost at least  $\epsilon$ , then the “effective depth” is roughly  $C^*/\epsilon$
  - Takes time  $O(b^{C^*/\epsilon})$  (exponential in effective depth)
- How much space does the fringe take?
  - Has roughly size of the last tier, so  $O(b^{C^*/\epsilon})$
- Is it complete?
  - Assuming best solution has a finite cost and minimum arc cost is positive, yes!
- Is it optimal?
  - Yes! (Proof soon via A\* algorithm)





# Performance Comparison

Algorithm		Complete	Optimal	Time	Space
DFS	w/ Path Checking	Y	N	$O(b^m)$	$O(bm)$
BFS		Y	Y*	$O(b^d)$	$O(b^d)$
UCS		Y*	Y	$O(b^{C^*/\epsilon})$	$O(b^{C^*/\epsilon})$



# Uniform Cost Issues


- Remember: UCS explores increasing cost contours
- The good: UCS is complete and optimal!
- The bad:
  - Explores options in every “direction”
  - No information about goal location
- We’ll fix that next!



# Uniform Cost: Pac-Man


---

- Cost of 1 for each action
- Explores all of the states, but one



# Search Heuristics


- Any estimate of how close a state is to a goal
- Designed for a particular search problem



- Examples: Manhattan distance, Euclidean distance

<https://qiao.github.io/PathFinding.js/visual/>

# Example: Heuristic Function




Straight-line distance  
to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

$h(x)$

# Example: Heuristic Function



Straight-line distance  
to Bucharest

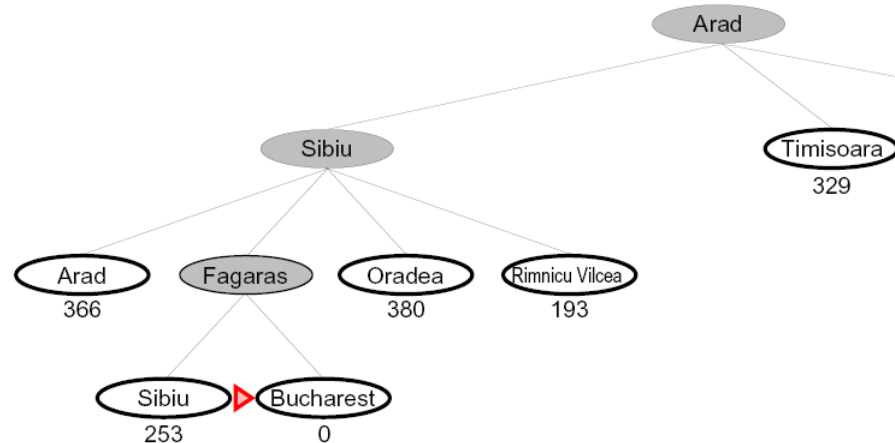
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

$h(x)$


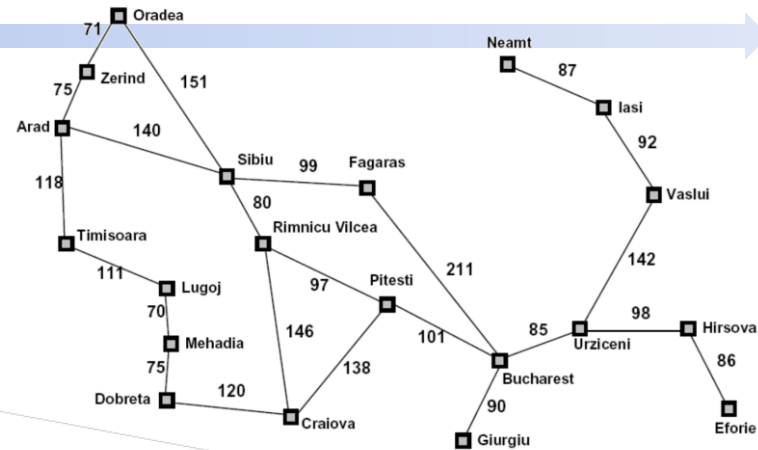


# Greedy Search

- Expand the node that seems closest...




- What can go wrong?



# Exercise: Greedy Search

- **Greedy** orders PQ by goal proximity, or *heuristic cost*  $h(n)$



- What is the greedy solution?



# Greedy ... in 3 lines 😊

```
def BFS(problem):  
    """Search the shallowest nodes in the search tree first."""  
    return greedy_search(problem,  
                           util.PriorityQueueWithFunction(heuristic))
```

```
def heuristic(Node n):  
    return 42
```


# Greedy Search (analysis sketch)

- Strategy: expand a node that you think is closest to a goal state
  - Heuristic: estimate of distance to nearest goal for each state
- A common case:
  - Best-first takes you straight to the (wrong) goal
- Worst-case: like a badly-guided DFS




# Video of Demo Contours Greedy (Empty)

---



# Video of Demo Contours Greedy (Pacman Small Maze)

---




# A\* Search



# Can we Combine UCS and Greedy?

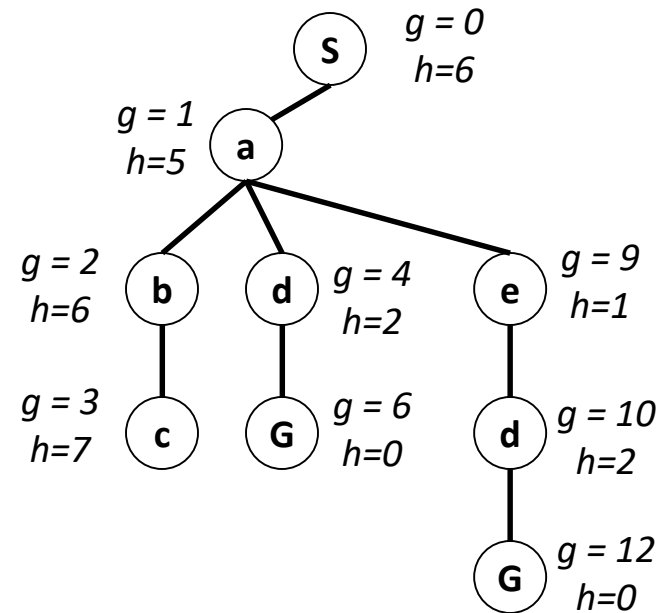

- **Uniform-cost** orders by path cost, or *backward cost*  $g(n)$
- **Greedy** orders by goal proximity, or *forward cost*  $h(n)$





# UCS + Greedy = A\*


- **Uniform-cost** orders by path cost, or *backward cost*  $g(n)$
- **Greedy** orders by goal proximity, or *forward cost*  $h(n)$



**A\* Search** orders by the sum:  $f(n) = g(n) + h(n)$

Example: Teg Grenager


# A\* search example



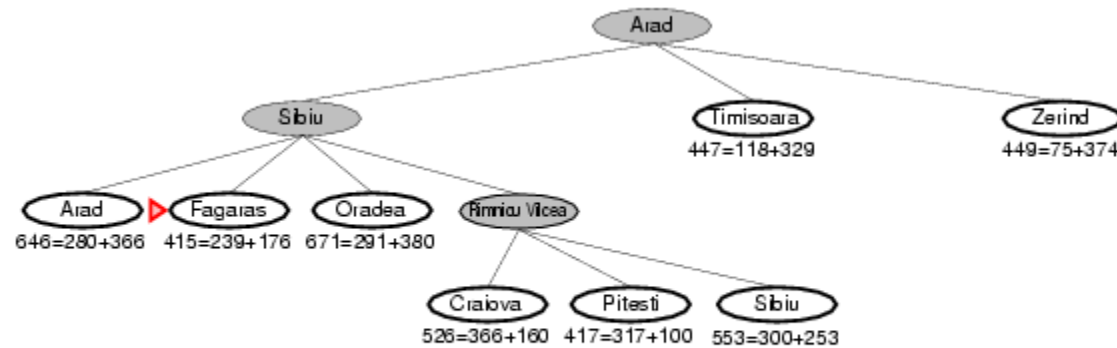
# A\* search example



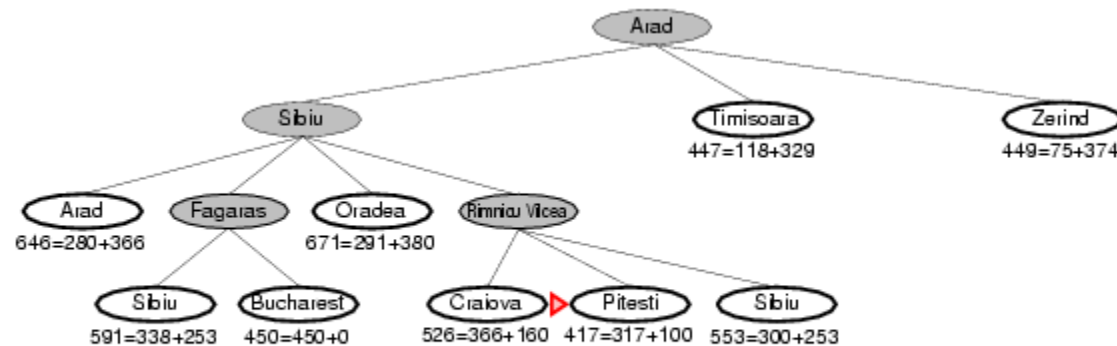
# A\* search example



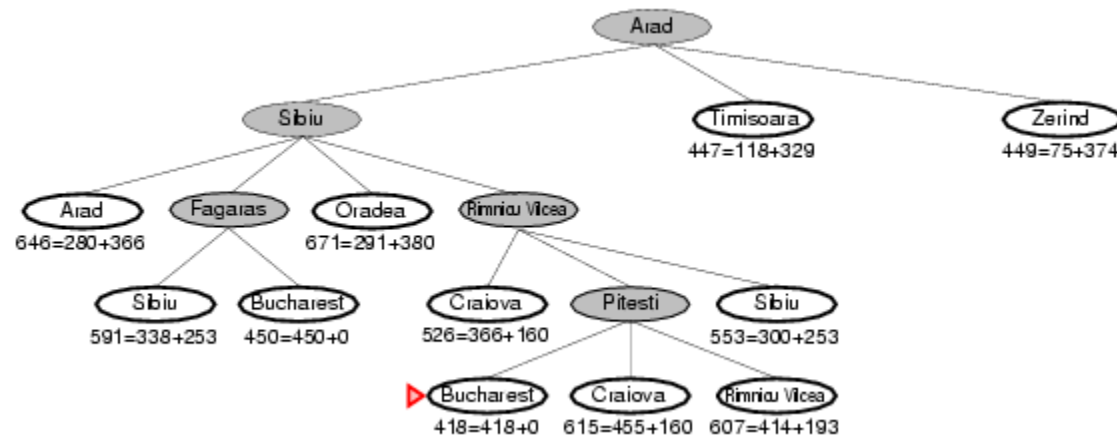
# A\* search example



# A\* search example



# A\* search example



# A\*... in 3 lines 😊

```
def BFS(problem):  
    """Search the shallowest nodes in the search tree first."""  
    return astar_search(problem,  
        util.PriorityQueueWithFunction(fcost))
```


```
def heuristic(Node n):  
    return ...
```

```
def fcost(Node n):  
    return heuristic(Node n) + n.getCost()
```




# When should $A^*$ terminate?


- Should we stop when we enqueue a goal?




## Exercise: A\* solution for this graph?



# Idea: Admissibility



Inadmissible (pessimistic) heuristics break optimality by trapping good plans on the fringe



Admissible (optimistic) heuristics slow down bad plans but never outweigh true costs


# Admissible Heuristics

- A heuristic  $h$  is *admissible* (optimistic) if:

$$0 \leq h(n) \leq h^*(n)$$

where  $h^*(n)$  is the true cost to a nearest goal

- Examples:



- Coming up with admissible heuristics is most of what's involved in using A\* in practice.


# Optimality of A\* Tree Search

Assume:

- A is an optimal goal node
- B is a suboptimal goal node
- $h$  is admissible

Claim:


- A will exit the fringe before B



# Optimality of A\* Tree Search: Blocking

Proof:

- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B
  1.  $f(n)$  is less or equal to  $f(A)$



$$f(n) = g(n) + h(n)$$

$$f(n) \leq g(A)$$

$$g(A) = f(A)$$

Definition of f-cost


Admissibility of  $h$

$h = 0$  at a goal

# Optimality of A\* Tree Search: Blocking

Proof:

- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B
  1.  $f(n)$  is less or equal to  $f(A)$
  2.  $f(A)$  is less than  $f(B)$



$$g(A) < g(B)$$

$$f(A) < f(B)$$


B is suboptimal

$h = 0$  at a goal

# Optimality of A\* Tree Search: Blocking

Proof:

- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B
  1.  $f(n)$  is less or equal to  $f(A)$
  2.  $f(A)$  is less than  $f(B)$
  3.  $n$  expands before B
- All ancestors of A expand before B
- A expands before B
- A\* search is optimal




$$f(n) \leq f(A) < f(B)$$




# Properties of A\*

Uniform-Cost





A\*




# UCS vs A\* Contours

- Uniform-cost expands equally in all “directions”
- A\* expands mainly toward the goal, but does hedge its bets to ensure optimality




# Video of Demo Contours (Empty) -- UCS

---




# Video of Demo Contours (Empty) -- Greedy

---




# Video of Demo Contours (Empty) – $A^*$


---




# Video of Demo Contours (Pacman Small Maze) – A\*



# Which Algorithm (1)?




# Which Algorithm (2)







# Which Algorithm (3)



# Which Algorithm (4)?



# Which Algorithm (5)



# Comparison: Summary



Greedy




Uniform Cost



A\*


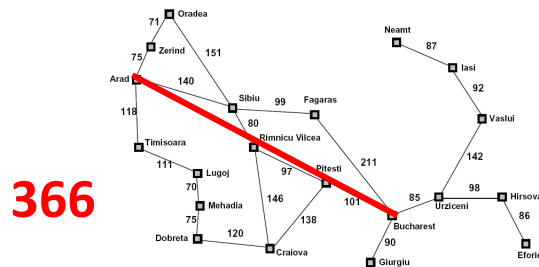
# A\* Applications

- Video games
- Pathing / routing problems
- Resource planning problems
- Robot motion planning
- Language analysis
- Machine translation
- Speech recognition
- ...



# Creating Admissible Heuristics

- Most of the work in solving hard search problems optimally is in coming up with admissible heuristics
- Often, admissible heuristics are solutions to *relaxed problems*, where new actions are available




- Inadmissible heuristics can be useful too!

# Example: 8 Puzzle

7	2	4
5		6
8	3	1

Start State



Actions

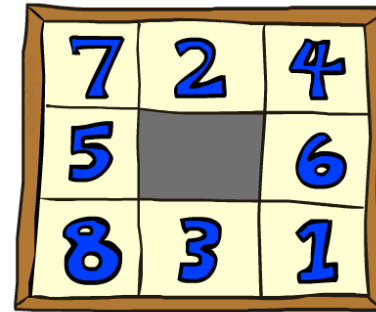
	1	2
3	4	5
6	7	8

Goal State

- What are the states?
- How many states?
- What are the actions?
- How many successors from the start state?
- What should the costs be?

# 8 Puzzle I

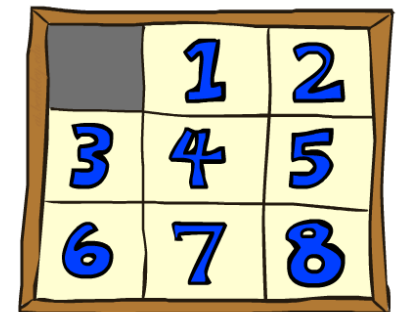
- Heuristic:
- Is it admissible?
- $h(\text{start}) =$
- $h(\text{goal}) =$



A 3x3 grid representing the start state of an 8-puzzle. The tiles are numbered 1 through 8, with the top-left cell being empty (gray). The numbers are in blue with black outlines.

	7	2	4
5			6
8	3	1	

Start State



A 3x3 grid representing the goal state of an 8-puzzle. The tiles are numbered 1 through 8, with the top-left cell being empty (gray). The numbers are in blue with black outlines.


	1	2
3	4	5
6	7	8

Goal State




# 8 Puzzle: Tiles heuristic


- Heuristic: **Number of tiles misplaced**
- Why is it admissible?
- $h(\text{start}) = 8$
- This is a *relaxed-problem* heuristic



Start State



Goal State




Average nodes expanded when the optimal path has...			
	...4 steps	...8 steps	...12 steps
UCS	112	6,300	$3.6 \times 10^6$
TILES	13	39	227


Statistics from Andrew Moore

# 8 Puzzle II: **Manhattan** heuristic

- **Relaxation**: easier 8-puzzle where **any tile could slide any direction at any time, ignoring other tiles?**
- Total **Manhattan** distance from correct **location**
- Is it admissible?
- $h(\text{start}) = 3 + 1 + 2 + \dots = 18$



Start State





Goal State

Average nodes expanded when the optimal path has...			
	...4 steps	...8 steps	...12 steps
TILES	13	39	227
MANHATTAN	12	25	73

# 8 Puzzle III: Oracle heuristic

- How about using the *actual cost* as a heuristic?
  - Would it be admissible?
  - Would we save on nodes expanded?
  - What's wrong with it?



- With A\*: a **trade-off** between **quality of estimate** and **work per node**
  - As heuristics get closer to the true cost, you will **expand fewer nodes** but usually **do more work per node** to compute the heuristic itself

# Recap: Problem Relaxation

- A problem with fewer restrictions on the actions is called a **relaxed problem**
- The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem
- If the rules of the 8-puzzle are relaxed so that a tile can move **anywhere**, then  $h_1(n)$  gives the shortest solution
- If the rules are relaxed so that a tile can move to **any adjacent square**, then  $h_2(n)$  gives the shortest solution

# Designing heuristics (cont'd)

E.g., for the 8-puzzle:

- $h_1(n)$  = number of misplaced tiles
- $h_2(n)$  = total Manhattan distance  
(i.e., no. of squares from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- $h_1(S) = ?$
- $h_2(S) = ?$
-

# Heuristics: cont'd

E.g., for the 8-puzzle:

- $h_1(n)$  = number of misplaced tiles
- $h_2(n)$  = total Manhattan distance  
(i.e., no. of squares from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- $h_1(S) = ?$  8
- $h_2(S) = ?$   $3+1+2+2+2+3+3+2 = 18$

Which is "better"  
–  $h_1$  or  $h_2$ ?

# Idea: Heuristic **dominance**

- If  $h_2(n) \geq h_1(n)$  for all  $n$  (both admissible, i.e.,  $<$  true cost)  
then  $h_2$  **dominates**  $h_1$

$h_2$  is better for search

- Typical search costs (average number of nodes expanded):
- $d=12$       IDS = 3,644,035 nodes  
     $A^*(h_1) = 227$  nodes  
     $A^*(h_2) = 73$  nodes
- $d=24$       IDS = too many nodes  
     $A^*(h_1) = 39,135$  nodes  
     $A^*(h_2) = 1,641$  nodes

# Example: Heuristics for Chess

- To select next move, must evaluate expected benefit of successor position:
  - **Value of the pieces** (count value of your pieces – value of opponents pieces)
  - **Space**: threatened/controlled space by you – space controlled by opponent
  - **Pawn** structure
  - ...
- Examples:
  - <https://www.quora.com/What-are-some-heuristics-for-quickly-evaluating-chess-positions>
  - <https://chessprogramming.wikispaces.com/Killer+Heuristic>



# Example: Heuristics for Motion Planning

- Robot motion: many moving (body) parts
- What's the most efficient way to accomplish goal?

<https://www.youtube.com/watch?v=dSwDZmvtGZY>

# Designing Heuristics


- A good heuristic is:
  - ✓ Admissible (optimistic)
  - Consistent (non-decreasing)
  - ✓ “Accurate”

# Trivial Heuristics, Dominance



- Dominance:  $h_a \geq h_c$  if
$$\forall n : h_a(n) \geq h_c(n)$$
- Max of admissible heuristics is admissible:

$$h(n) = \max(h_a(n), h_b(n))$$


- Trivial heuristics
  - Bottom of lattice is the zero heuristic (what does this give us?)
  - Top of lattice is the exact heuristic



# A\* Graph Search Gone Wrong?




# Consistency of Heuristics



- Main idea: estimated heuristic costs  $\leq$  actual costs
  - Admissibility: heuristic cost  $\leq$  actual cost to goal  
 $h(A) \leq \text{actual cost from A to G}$
  - Consistency: heuristic “arc” cost  $\leq$  actual cost for each arc  
 $h(A) - h(C) \leq \text{cost(A to C)}$
- Consequences of consistency:
  - The f value along a path never decreases  
 $h(A) \leq \text{cost(A to C)} + h(C)$
  - A\* graph search is optimal


# Optimality of A\* Graph Search

- Sketch of proof: consider what A\* does with a **consistent** heuristic:
  - **Fact 1:** In tree search, A\* expands nodes in increasing total  $f$  value ( $f$ -contours)
  - **Fact 2:** For every state  $s$ , nodes that reach  $s$  optimally are expanded before nodes that reach  $s$  suboptimally
  - Result: A\* graph search is optimal



# Optimality (2): Tree vs. Graph Search

- **Tree search:**
  - A\* is optimal if heuristic is **admissible**
  - UCS is a special case ( $h = 0$ )
- **Graph search:**
  - A\* optimal if heuristic is **consistent**
  - UCS optimal ( $h = 0$  is consistent)
- Consistency implies admissibility
- In general, **most natural admissible heuristics tend to be consistent**, especially if from relaxed problems




# A\*: Summary





# A\*: Summary

- A\* uses both backward costs and (estimates of) forward costs
- A\* is optimal with admissible / consistent heuristics
- Heuristic design is key: often use relaxed problems



# Project 1: Due **Friday, Feb 9**

- Read FAQ on Canvas before posting questions:
- **Questions 1-4:** if you develop a correct solution for DFS, the rest will be easy modifications
- **Run autograder** after \*every\* question. Until you perfectly pass all the test cases, assume your code has bugs.
- Example (incomplete!) implementations:  
<https://github.com/aimacode/aima-python/blob/master/search.py>

# Tips for Project 1 (cont'd)

- Problems 5-8 depend on code in 1-4. Get that right (and tested) first, before moving on!
- P5/Corners problem: must visit all corners in \*single\* path
  - Implications for search tree, state info to update
- Heuristics for p6-8: start simple. For extra credit, think back to graph traversal algorithms from cs323.

# Project 1 final hints

- Use Discussions, read FAQ before posting questions
- **Suggestion: use Node class or similar.** Easier to do Questions 5-8.
- Questions 5-8: more fun/creative. Leave enough time, start early.
- **Most importantly:** Don't Panic! Eat the elephant one question at a time.