

(Hidden) Markov Models

Wrap-up

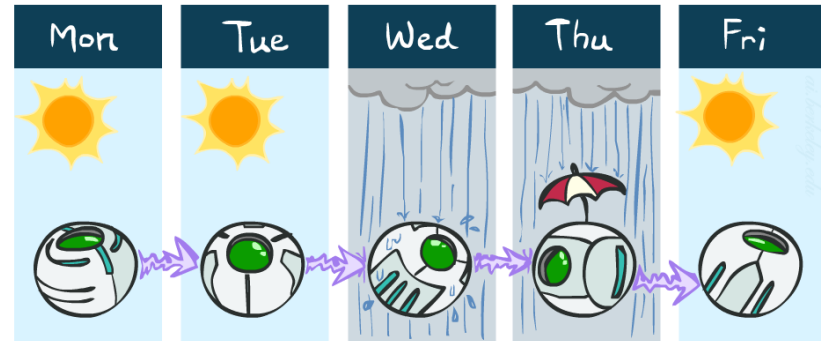
With slides from Dan Klein and Pieter Abbeel

Today

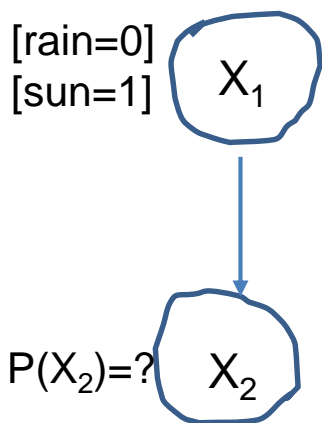
- Review Project 2 solution (optional)
- BN Inference with time
 - Markov Chains
 - Hidden Markov Models
- Project 3: Ghostbusters
- Wrap-up: Search and Inference for “intelligent” agents
- **[optional] Midterm review today 5:30pm in W301**

Example Markov Chain: Weather

- States: $X = \{\text{rain}, \text{sun}\}$
- Initial distribution: 1.0 sun

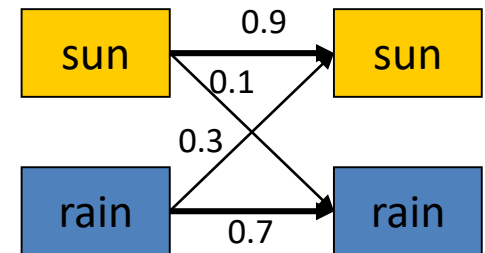
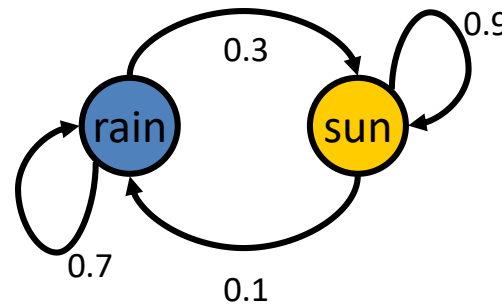


Two new ways of representing the same BN



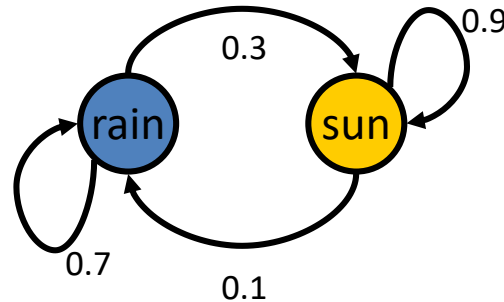
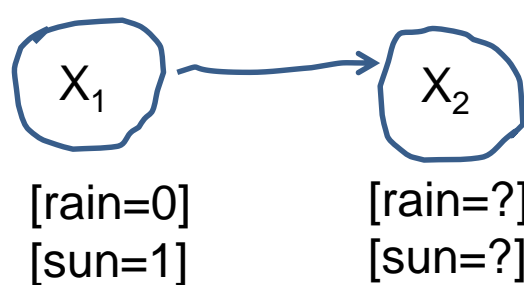
CPT $P(X_t | X_{t-1})$:

X_{t-1}	X_t	$P(X_t X_{t-1})$
sun	sun	0.9
sun	rain	0.1
rain	sun	0.3
rain	rain	0.7



Weather Markov Chain: Prediction

- Initial distribution: 1.0 sun



X_{t-1}	X_t	$P(X_t X_{t-1})$
sun	sun	0.9
sun	rain	0.1
rain	sun	0.3
rain	rain	0.7

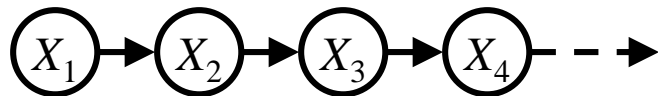
- What is the probability distribution after one step?

$$P(X_2 = \text{sun}) = P(X_2 = \text{sun} | X_1 = \text{sun})P(X_1 = \text{sun}) + P(X_2 = \text{sun} | X_1 = \text{rain})P(X_1 = \text{rain})$$

$$0.9 \cdot 1.0 + 0.3 \cdot 0.0 = 0.9$$

Prediction: Mini-Forward Algorithm

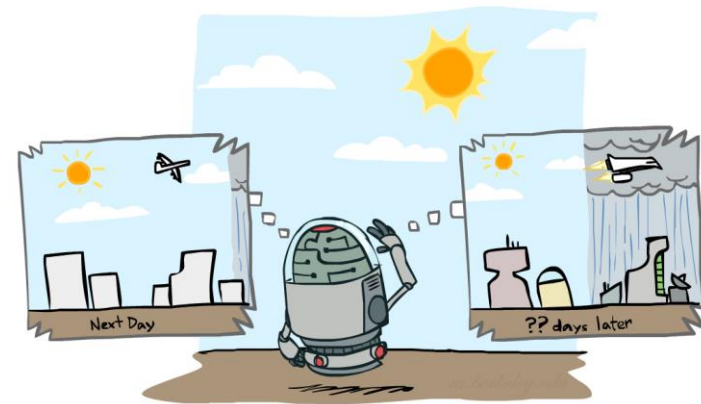
- Question: What's $P(X)$ on some day t ?



$$P(x_1) = \text{known}$$

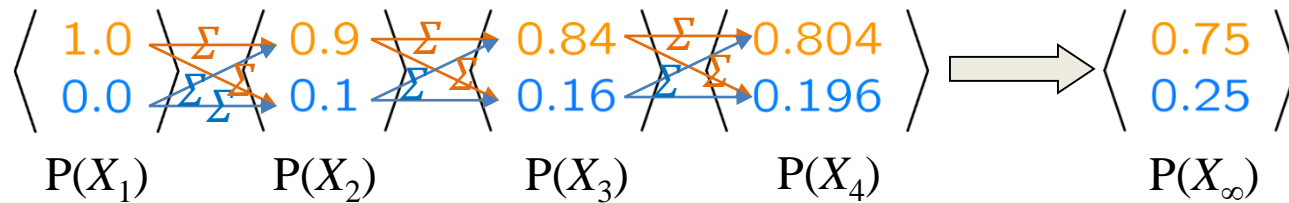
$$\begin{aligned} P(x_t) &= \sum_{x_{t-1}} P(x_{t-1}, x_t) \\ &= \sum_{x_{t-1}} P(x_t \mid x_{t-1}) P(x_{t-1}) \end{aligned}$$

Forward simulation

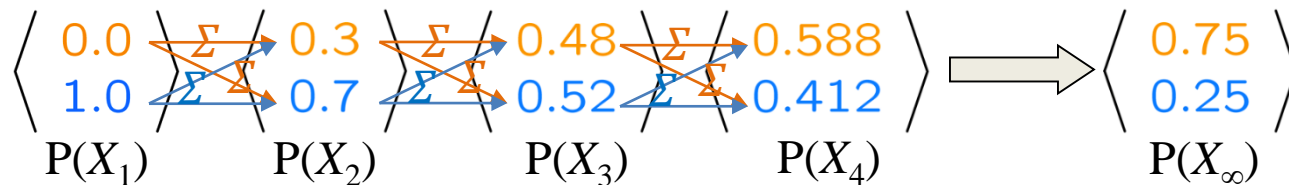


Example Run of Mini-Forward Algorithm

- From initial observation of sun



- From initial observation of rain



- From yet another initial distribution $P(X_1)$:



Markov Chains: (Simplified)? Matrix View

- A probability vector $\mathbf{x} = (x_1, \dots, x_n)$ tells us our *belief distribution which state we are at a given time*.
- For convenience we represent X as a *row vector*
 $\mathbf{X}_1: [1, 0]$ (sun=1, rain=0)
 $\mathbf{X}_2: [0.9, 0.1]$ (sun=0.9, rain=0.1)
- Transition Probability *Matrix* $\mathbf{P} =$ (re-formatted) CPT

CPT $P(X_t | X_{t-1})$:

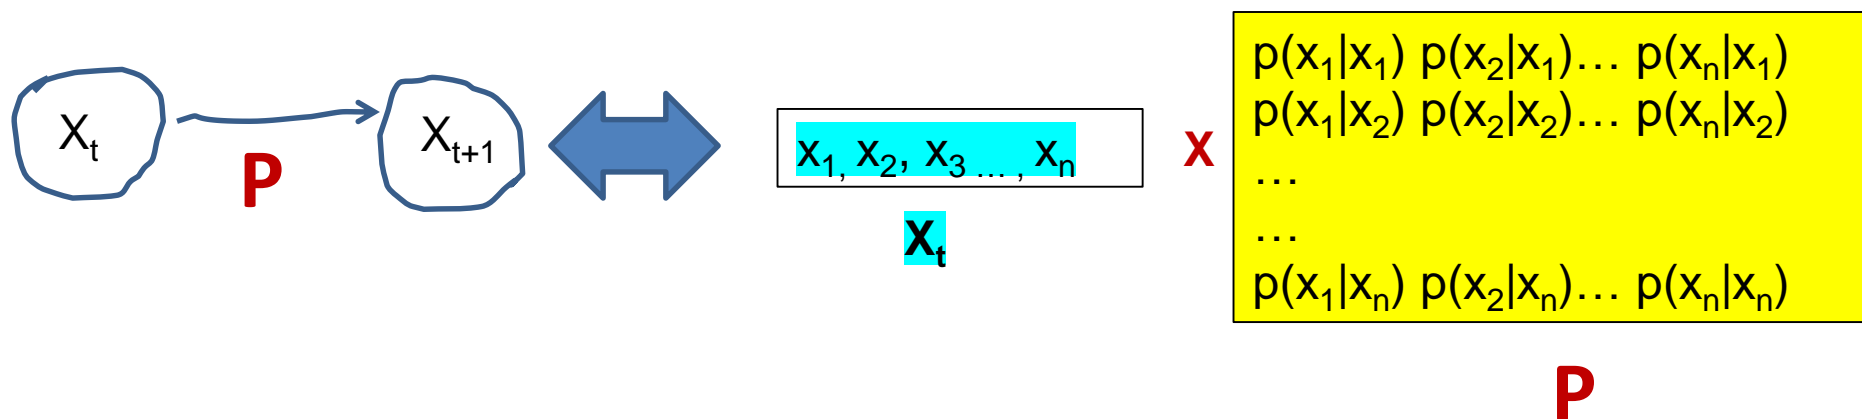
X_{t-1}	X_t	$P(X_t X_{t-1})$
sun	sun	0.9
sun	rain	0.1
rain	sun	0.3
rain	rain	0.7



$X_t \backslash X_{t+1}$	sun	rain
sun	0.9	0.1
rain	0.3	0.7

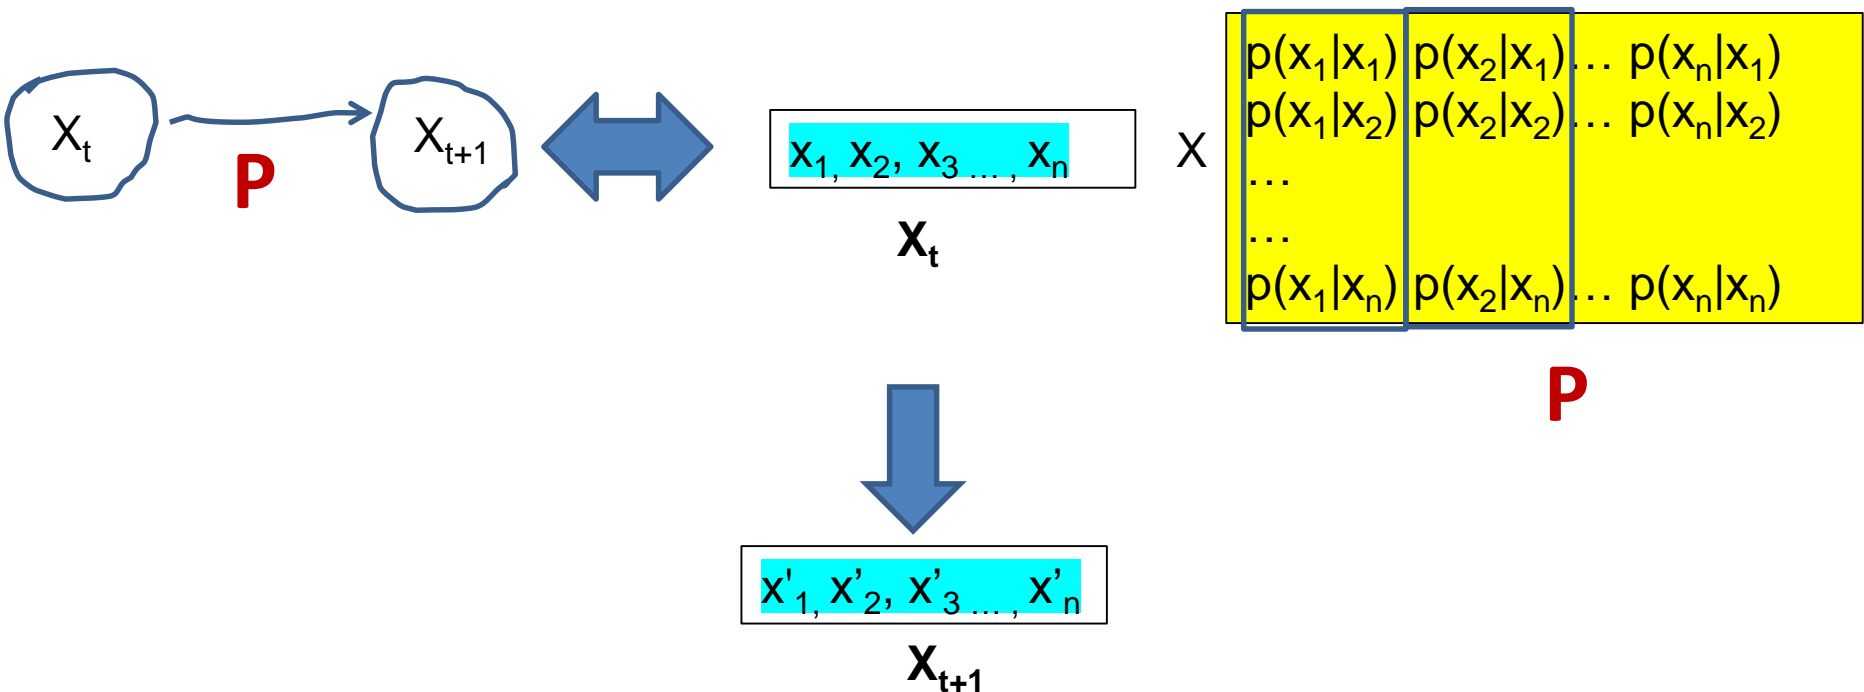
Forward Algorithm in vector form

- If the probability vector is $\mathbf{x} = (x_1, \dots, x_n)$ at this step, what is \mathbf{x}' at the next step?
- Recall that row i of **the transition probability Matrix \mathbf{P}** tells us where we go next from state i .
- So from \mathbf{x} , our next state is computed as $\mathbf{x}\mathbf{P}$.



Forward Algorithm in vector form

- So from \mathbf{x} , our next belief/state Prob is \mathbf{xP} .



$$x'_1 = x_1 * p(x_1|x_1) + x_2 * p(x_1|x_2) + x_3 * p(x_1|x_3) + \dots + x_n * p(x_1|x_n)$$

$$x'_2 = x_1 * p(x_2|x_1) + x_2 * p(x_2|x_2) + x_3 * p(x_2|x_3) + \dots + x_n * p(x_2|x_n)$$

Stationary Distributions

- For most chains:

- Influence of the initial distribution gets less and less over time.
- The distribution we end up in is independent of the initial distribution

- **Stationary distribution:**

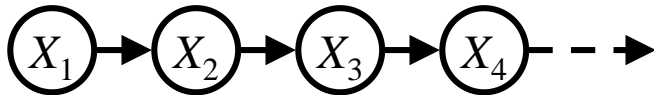
- The distribution we end up with is called the **stationary distribution** P_∞ of the Markov chain
- It satisfies

$$P_\infty(X) = P_{\infty+1}(X) = \sum_x P(X|x)P_\infty(x)$$



Example: Stationary Distributions

- Question: What's $P(X)$ at time $t = \text{infinity}$?



1 $P_{\infty}(\text{sun}) = P(\text{sun}|\text{sun})P_{\infty}(\text{sun}) + P(\text{sun}|\text{rain})P_{\infty}(\text{rain})$
 $P_{\infty}(\text{rain}) = P(\text{rain}|\text{sun})P_{\infty}(\text{sun}) + P(\text{rain}|\text{rain})P_{\infty}(\text{rain})$

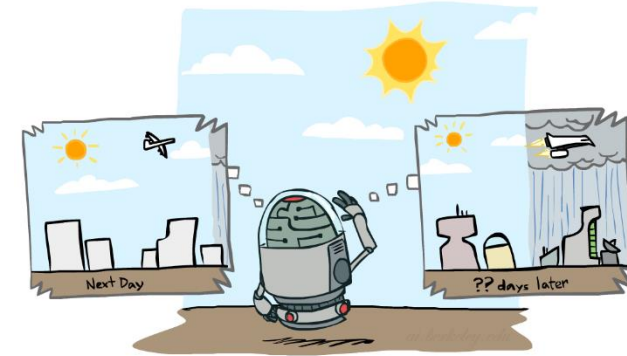
2 $P_{\infty}(\text{sun}) = 0.9P_{\infty}(\text{sun}) + 0.3P_{\infty}(\text{rain})$
 $P_{\infty}(\text{rain}) = 0.1P_{\infty}(\text{sun}) + 0.7P_{\infty}(\text{rain})$

3 $P_{\infty}(\text{sun}) = 3P_{\infty}(\text{rain})$
 $P_{\infty}(\text{rain}) = 1/3P_{\infty}(\text{sun})$

4 Also: $P_{\infty}(\text{sun}) + P_{\infty}(\text{rain}) = 1$

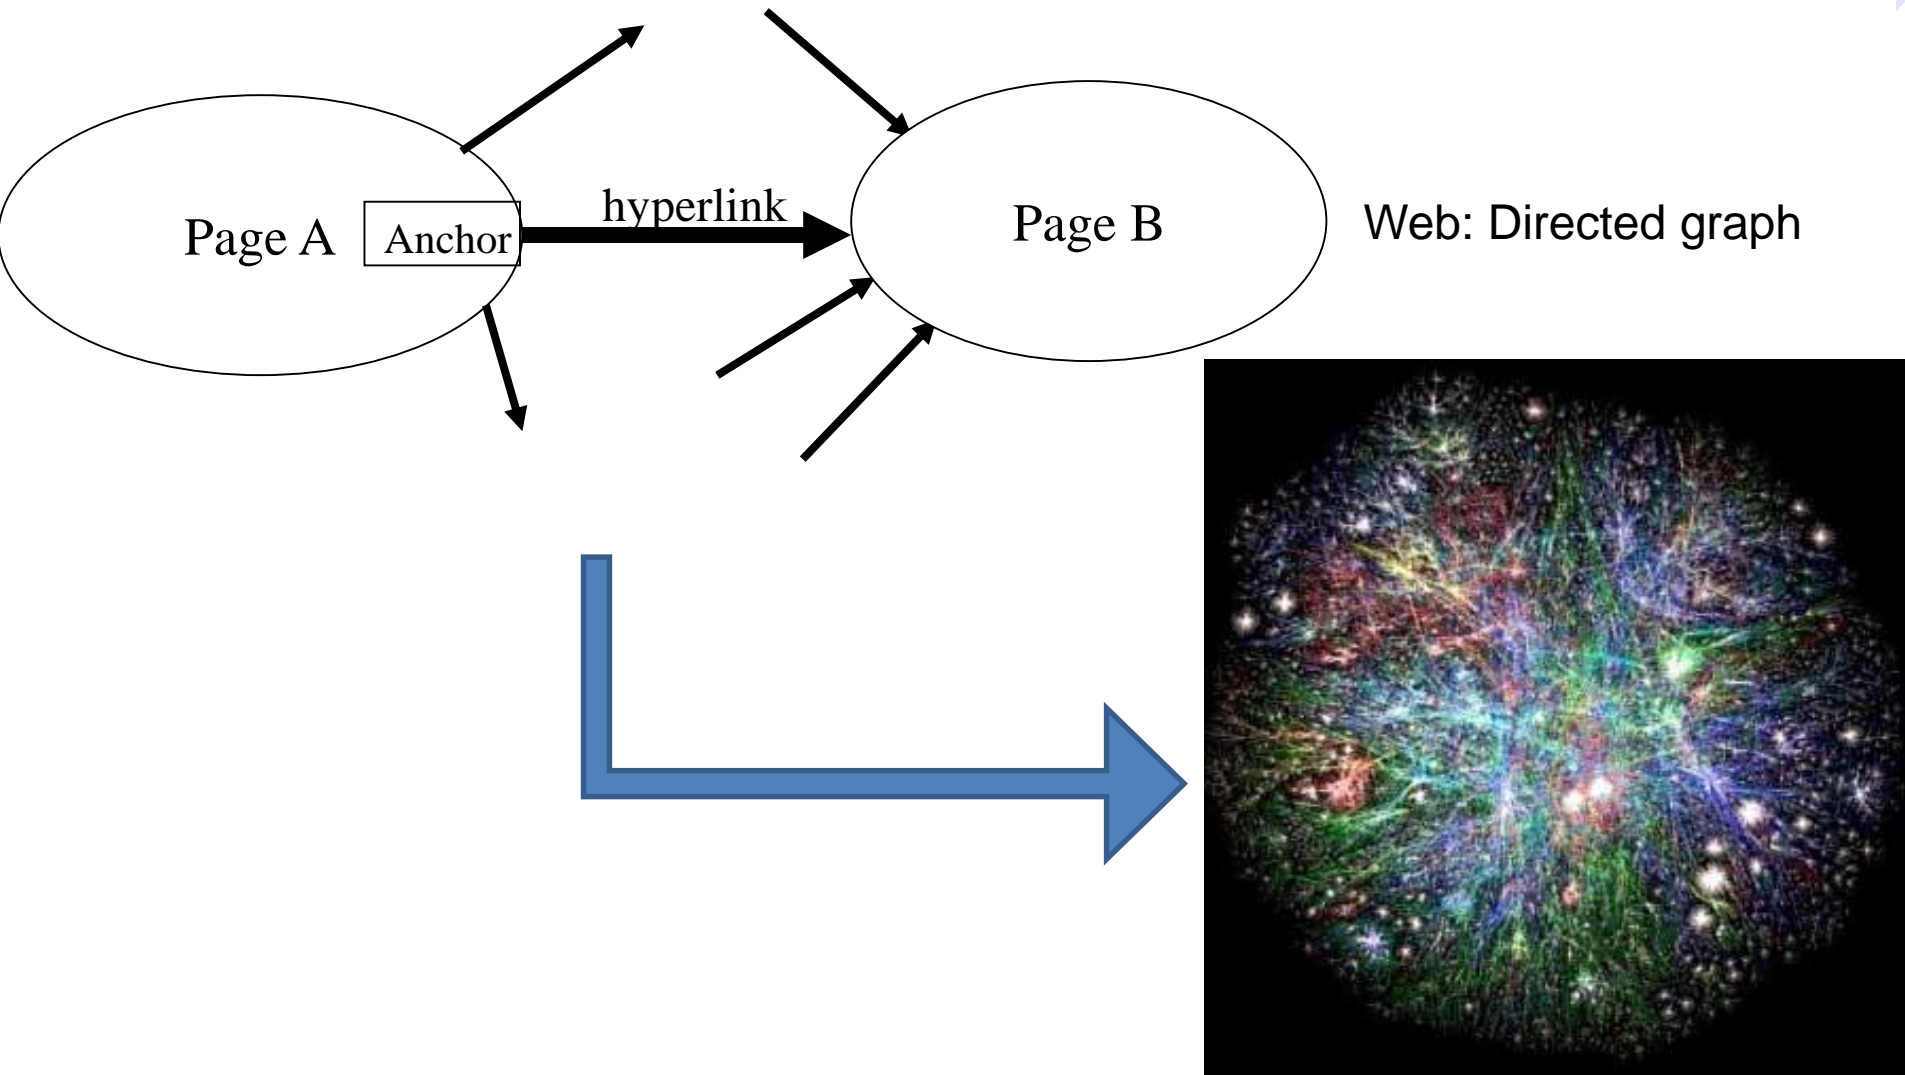


$$P_{\infty}(\text{sun}) = 3/4$$
$$P_{\infty}(\text{rain}) = 1/4$$

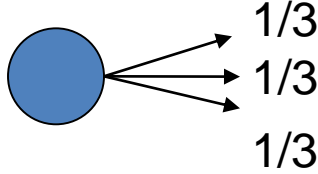


X_{t-1}	X_t	$P(X_t X_{t-1})$
sun	sun	0.9
sun	rain	0.1
rain	sun	0.3
rain	rain	0.7

Application: PageRank

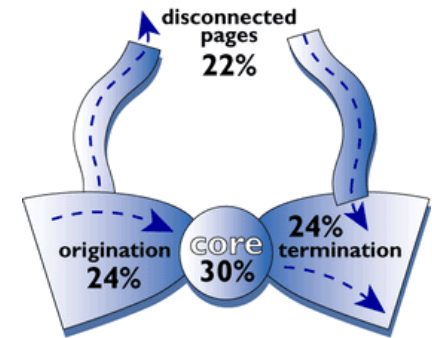


Random (Markov) Surfer Model

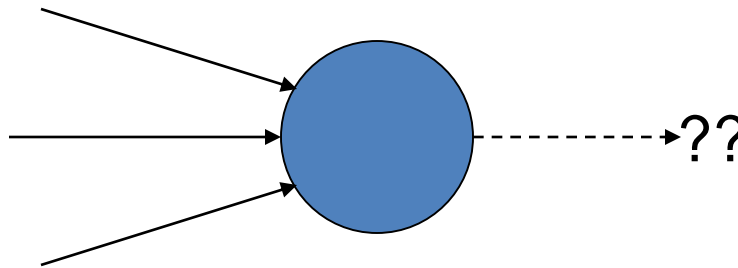
- Imagine a browser doing a random walk on web pages:
 - Start at a random page
 - A blue circular node with three arrows pointing to the right. Each arrow is labeled with the fraction $1/3$.
 - At each step, go out of the current page along one of the links on that page, with uniform prob.
- “In the steady state” each page has a long-term visit rate - use this as the page’s PageRank score.
 - Only works for “ergodic” chains (no dead-ends)

Not quite enough

- The web is full of dead-ends.

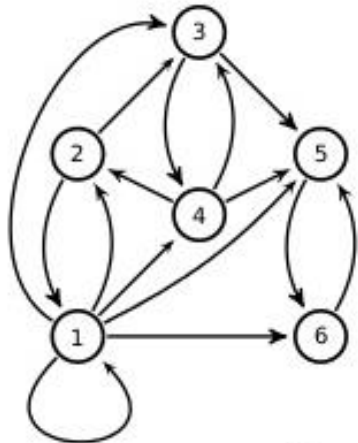


- Random walk can get stuck in dead-ends.
- Makes no sense to talk about long-term visit rates.



Add Teleportation

- Introduce a “random teleportation” probability, $E(p)$, that sends a surfer to any page with small (random) probability.



The random surfer model

At a node ...

1. follow edges with prob α
2. do something else with prob $(1 - \alpha)$

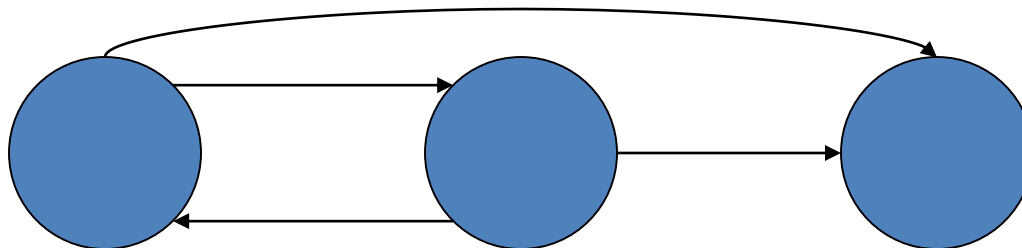


The important pages are the places we are most likely to find the random surfer

$$R(p) = c \left(\sum_{q:q \rightarrow p} \frac{R(q)}{N_q} + E(p) \right)$$

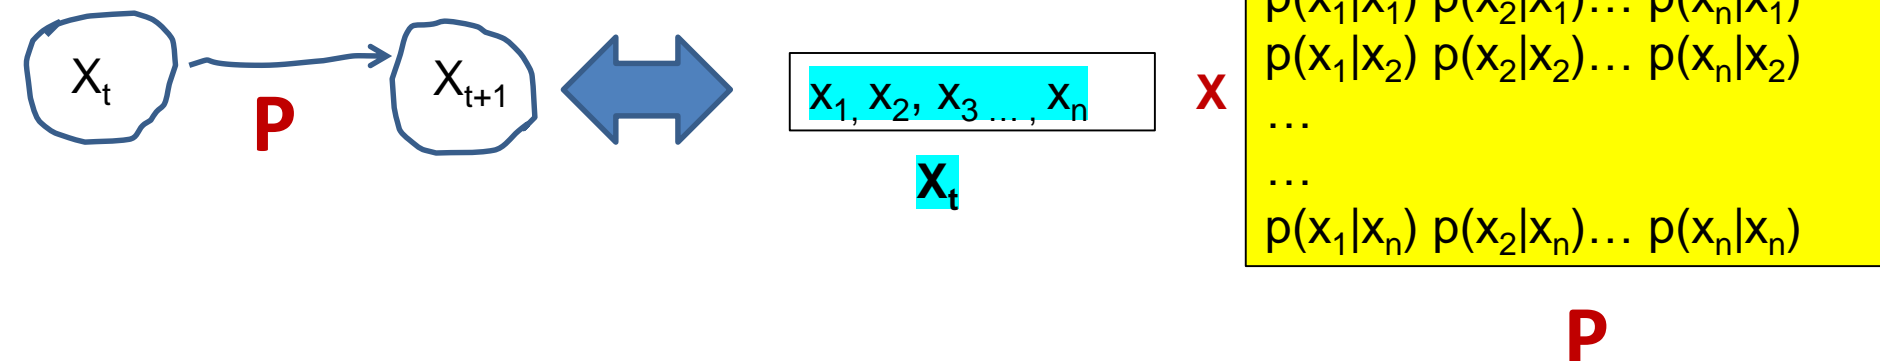
$R(p)$ = Steady State Markov Prob!

- For all pages i , $\sum_{j=1}^n P_{ij} = 1$. (normalization)
- Represent the teleporting random walk as a Markov chain, for this case:



Iterative (Forward) Way of computing \mathbf{a}

- Recall, regardless of where we start, we eventually reach the steady state \mathbf{a} .
- Start with any distribution (say $\mathbf{x}=(10...0)$).
- After one step, we're at \mathbf{xP} ;
- after two steps at \mathbf{xP}^2 , then \mathbf{xP}^3 and so on.
- “Eventually” means for “large” k , $\mathbf{xP}^k = \mathbf{a}$.
- Algorithm: multiply \mathbf{x} by increasing powers of \mathbf{P} until the product converges.



*PageRank Algorithm



*Page = *Larry Page*, not web page!

Let S be the total set of pages.

Let $\forall p \in S: E(p) = \alpha/|S|$ (for some $0 < \alpha < 1$, e.g. 0.15)

Initialize $\forall p \in S: R(p) = 1/|S|$

Until ranks do not change (much) (*convergence*)

For each $p \in S$:

$$R'(p) = \left[(1 - \alpha) \sum_{q: q \rightarrow p} \frac{R(q)}{N_q} \right] + E(p)$$

$$c = 1 / \sum_{p \in S} R'(p)$$

For each $p \in S: R(p) = cR'(p)$ (*normalize*)

Speed of Convergence

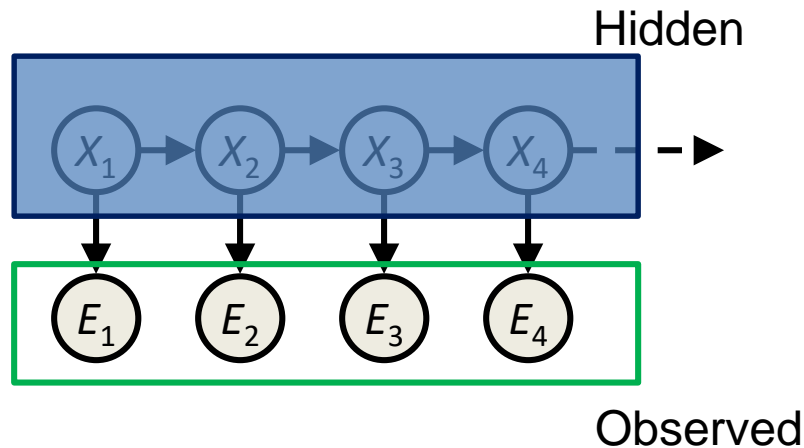
- Early experiments on Google used 322 million links.
- PageRank algorithm converged (within small tolerance) in about 52 iterations.
- Number of iterations required for convergence is empirically $O(\log n)$ (where n is the number of links).
- Speed of convergence shown related to sizes of eigenvectors of the matrix P

Which page has highest PageRank?

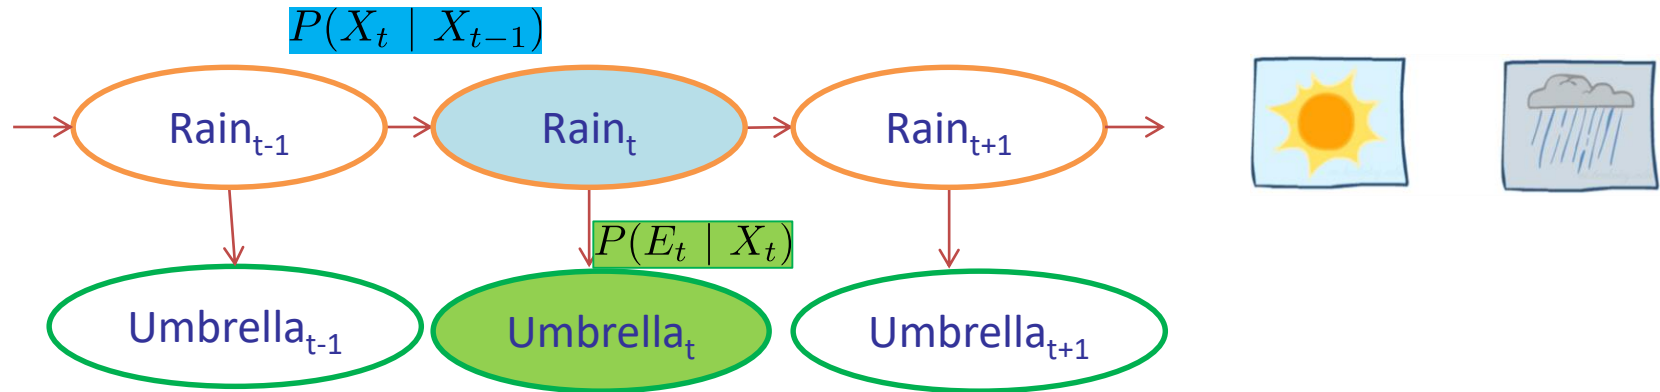
- c.f. 1997: Netscape!
- c.f. 2005: Wikipedia!
 - Maybe not (hard to measure externally)
- Some other sites with high PageRank
 - Google
 - MS Internet Explorer, Firefox homepages

Hidden Markov Models

- Markov chains not so useful for most agents
 - Need **observations** to update your beliefs
- Hidden Markov models (HMMs)
 - Underlying Markov chain over states X
 - **You observe outputs (effects) at each time step**



Example: Weather HMM

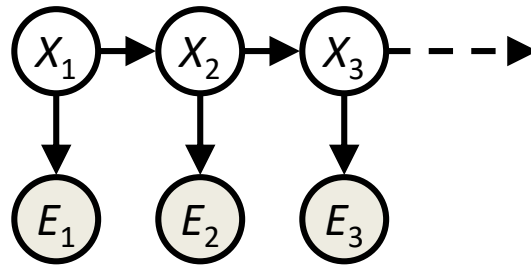


- An HMM is defined by:
 - Initial distribution: $P(X_1)$
 - Transitions: $P(X_t | X_{t-1})$
 - Emissions/observations: $P(E_t | X_t)$

R_t	R_{t+1}	$P(R_{t+1} R_t)$
+r	+r	0.7
+r	-r	0.3
-r	+r	0.3
-r	-r	0.7

R_t	U_t	$P(U_t R_t)$
+r	+u	0.9
+r	-u	0.1
-r	+u	0.2
-r	-u	0.8

Joint Distribution of an HMM



– Joint distribution:

$$P(X_1, E_1, X_2, E_2, X_3, E_3) = P(X_1)P(E_1|X_1)P(X_2|X_1)P(E_2|X_2)P(X_3|X_2)P(E_3|X_3)$$

– More generally:

$$P(X_1, E_1, \dots, X_T, E_T) = P(X_1)P(E_1|X_1) \prod_{t=2}^T P(X_t|X_{t-1})P(E_t|X_t)$$

– Questions to be resolved:

- Does this indeed define a joint distribution?
- Can every joint distribution be factored this way, or are we making some assumptions about the joint distribution by using this factorization?

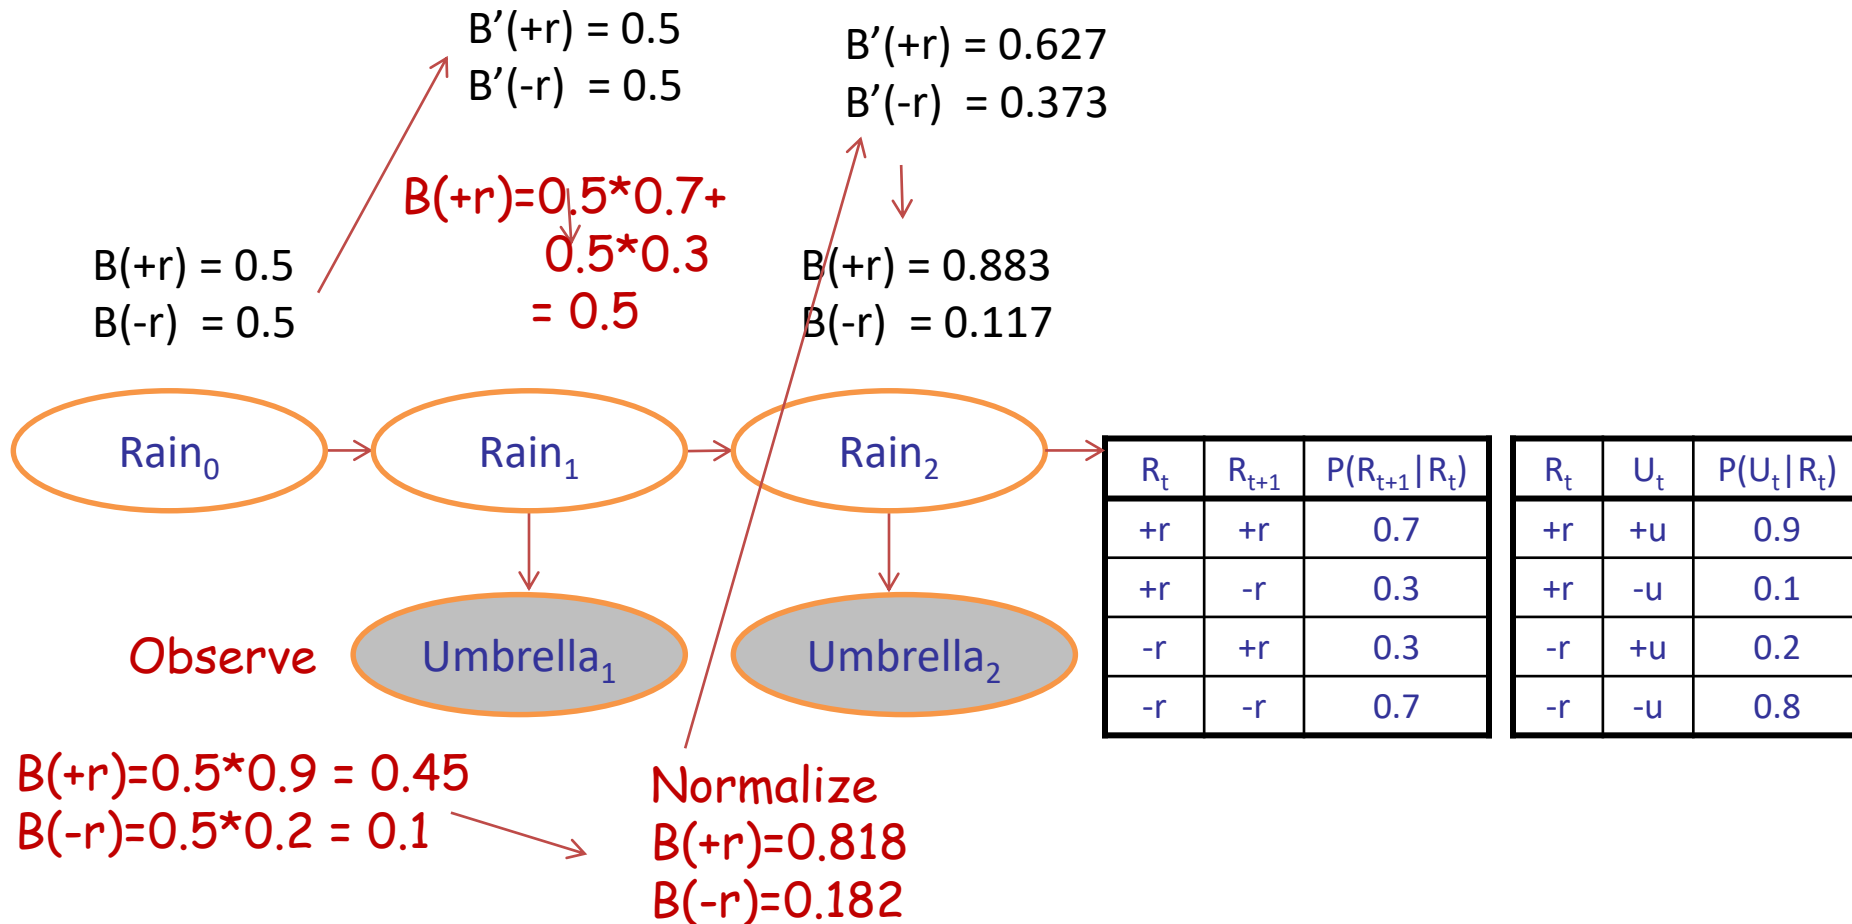
Filtering / Monitoring

- Filtering, or monitoring, is the task of tracking the distribution $B_t(X) = P_t(X_t \mid e_1, \dots, e_t)$ (the belief state) over time
- We start with $B_1(X)$ in an initial setting, usually uniform
- As time passes, or we get observations, we update $B(X)$

Filtering: Weather HMM



Marginalize $B(+r) = B'(+r) * P(+r|+r) + B'(-r) * P(+r|-r)$



Ghostbusters Filtering (project 3)

- Let's say we have two distributions:
 - **Prior distribution** over ghost location: $P(G)$
 - Let's say this is uniform
 - **Sensor reading model**: $P(R | G)$
 - Given: we know what our sensors do
 - R = reading color measured at (1,1)
 - E.g. $P(R = \text{yellow} | G=(1,1)) = 0.1$
- Can calculate **posterior distribution** $P(G|r)$ over ghost locations given a sensor reading, with Bayes' rule:

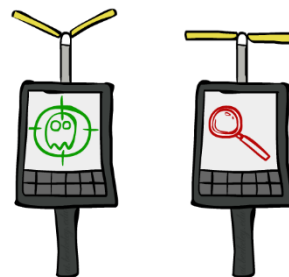
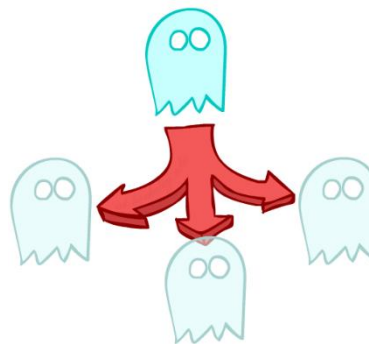
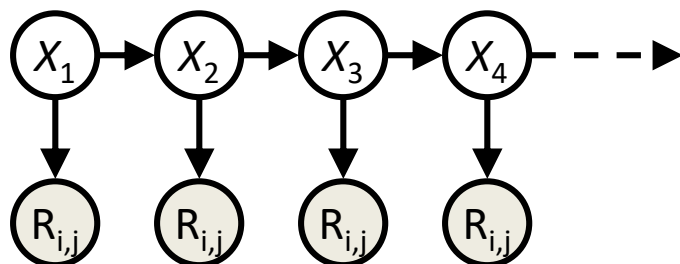
$$P(g|r) \propto P(r|g)P(g)$$

0.11	0.11	0.11
0.11	0.11	0.11
0.11	0.11	0.11

0.17	0.10	0.10
0.09	0.17	0.10
<0.01	0.09	0.17

Example: Ghostbusters HMM

- $P(X_1) = \text{uniform}$
- $P(X|X') = \text{usually move clockwise, but sometimes move in a random direction or stay in place}$
- $P(R_{ij}|X) = \text{same sensor model as before: red means close, green means far away.}$



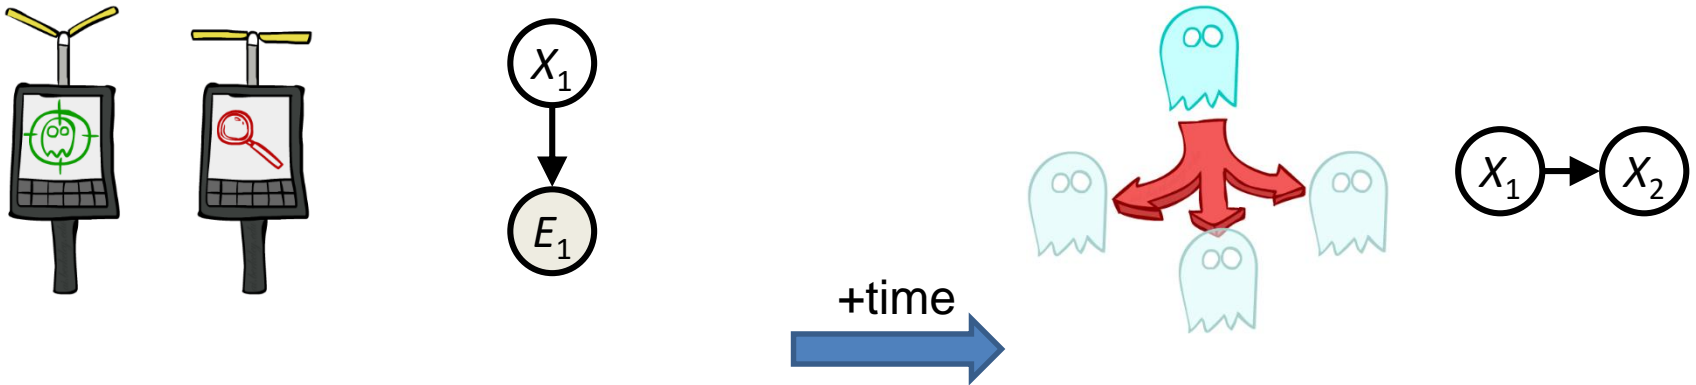
1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

$P(X_1)$

1/6	1/6	1/2
0	1/6	0
0	0	0

$P(X|X' = \langle 1, 2 \rangle)$

Ghost Localization: Base Cases (time=1)



$$P(X_1|e_1)$$

$$\begin{aligned} P(x_1|e_1) &= P(x_1, e_1)/P(e_1) \\ &\propto_{X_1} P(x_1, e_1) \\ &= P(x_1)P(e_1|x_1) \end{aligned}$$

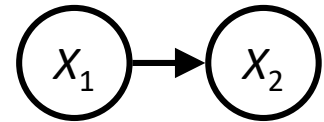
$$P(X_2)$$

$$\begin{aligned} P(x_2) &= \sum_{x_1} P(x_1, x_2) \\ &= \sum_{x_1} P(x_1)P(x_2|x_1) \end{aligned}$$

Ghost Localization: Time = 2

- Assume we have current belief $P(X \mid \text{evidence before now})$

$$B(X_t) = P(X_t | e_{1:t})$$



- Then, after one time step passes:

$$\begin{aligned} P(X_{t+1} | e_{1:t}) &= \sum_{x_t} P(X_{t+1}, x_t | e_{1:t}) \\ &= \sum_{x_t} P(X_{t+1} | x_t, e_{1:t}) P(x_t | e_{1:t}) \\ &= \sum_{x_t} P(X_{t+1} | x_t) P(x_t | e_{1:t}) \end{aligned}$$

- Or compactly:

$$B'(X_{t+1}) = \sum_{x_t} P(X' | x_t) B(x_t)$$

- Basic idea: beliefs get “pushed” through the transitions
 - With the “B” notation, we have to be careful about what time step t the belief is about, and what evidence it includes

Example: Passage of Time

- As time passes, uncertainty “accumulates”

(Transition model: ghosts usually go clockwise)

<0.01	<0.01	<0.01	<0.01	<0.01	<0.01
<0.01	<0.01	<0.01	<0.01	<0.01	<0.01
<0.01	<0.01	1.00	<0.01	<0.01	<0.01
<0.01	<0.01	<0.01	<0.01	<0.01	<0.01

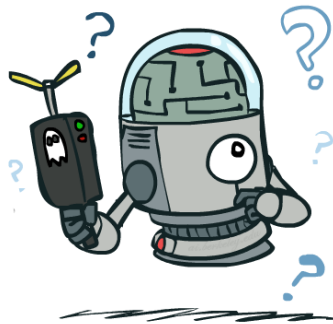
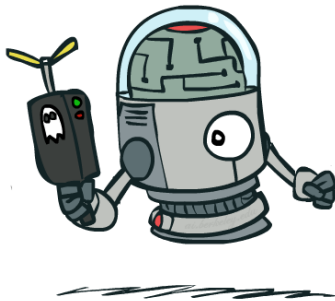
T = 1

<0.01	<0.01	<0.01	<0.01	<0.01	<0.01
<0.01	<0.01	0.06	<0.01	<0.01	<0.01
<0.01	0.76	0.06	0.06	<0.01	<0.01
<0.01	<0.01	0.06	<0.01	<0.01	<0.01

T = 2

0.05	0.01	0.05	<0.01	<0.01	<0.01
0.02	0.14	0.11	0.35	<0.01	<0.01
0.07	0.03	0.05	<0.01	0.03	<0.01
0.03	0.03	<0.01	<0.01	<0.01	<0.01

T = 5



Observation

- Assume we have current belief $P(X \mid \text{previous evidence})$:

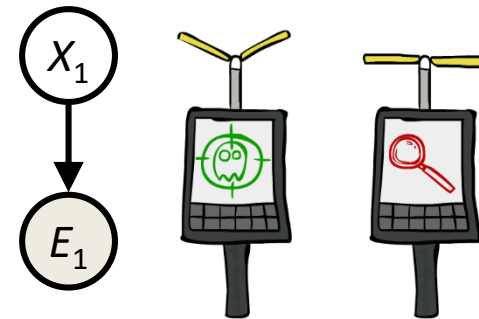
$$B'(X_{t+1}) = P(X_{t+1} | e_{1:t})$$

- Then, after evidence comes in:

$$\begin{aligned} P(X_{t+1} | e_{1:t+1}) &= P(X_{t+1}, e_{t+1} | e_{1:t}) / P(e_{t+1} | e_{1:t}) \\ &\propto_{X_{t+1}} P(X_{t+1}, e_{t+1} | e_{1:t}) \\ &= P(e_{t+1} | e_{1:t}, X_{t+1}) P(X_{t+1} | e_{1:t}) \\ &= P(e_{t+1} | X_{t+1}) P(X_{t+1} | e_{1:t}) \end{aligned}$$

- Or, compactly:

$$B(X_{t+1}) \propto_{X_{t+1}} P(e_{t+1} | X_{t+1}) B'(X_{t+1})$$



- Basic idea: beliefs “reweighted” by likelihood of evidence
- Unlike passage of time, we have to renormalize

Example: Observation

- As we get observations, beliefs get reweighted, uncertainty “decreases”

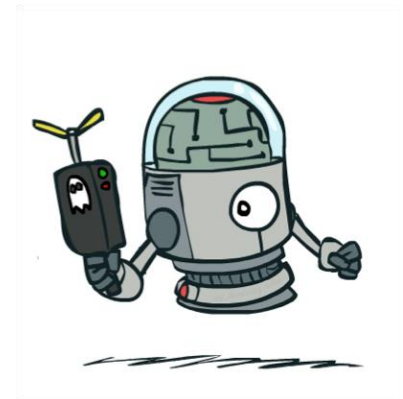
0.05	0.01	0.05	<0.01	<0.01	<0.01
0.02	0.14	0.11	0.35	<0.01	<0.01
0.07	0.03	0.05	<0.01	0.03	<0.01
0.03	0.03	<0.01	<0.01	<0.01	<0.01

Before observation

<0.01	<0.01	<0.01	<0.01	0.02	<0.01
<0.01	<0.01	<0.01	0.83	0.02	<0.01
<0.01	<0.01	0.11	<0.01	<0.01	<0.01
<0.01	<0.01	<0.01	<0.01	<0.01	<0.01

After observation

$$B(X) \propto P(e|X)B'(X)$$



HMM Inference: Forward Algorithm

- We are given evidence at each time and want to know

$$B_t(X) = P(X_t|e_{1:t})$$

- We can derive the following updates

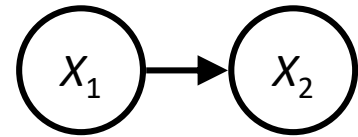
$$\begin{aligned} P(x_t|e_{1:t}) &\propto_X P(x_t, e_{1:t}) \\ &= \sum_{x_{t-1}} P(x_{t-1}, x_t, e_{1:t}) \\ &= \sum_{x_{t-1}} P(x_{t-1}, e_{1:t-1}) P(x_t|x_{t-1}) P(e_t|x_t) \\ &= P(e_t|x_t) \sum_{x_{t-1}} P(x_t|x_{t-1}) P(x_{t-1}, e_{1:t-1}) \\ &= P(e_t|x_t) \sum_{x_{t-1}} P(x_t|x_{t-1}) B'(x_{t-1}) \end{aligned}$$

We can normalize as we go if we want to have $P(x|e)$ at each time step, or just once at the end...

Online Belief Updates

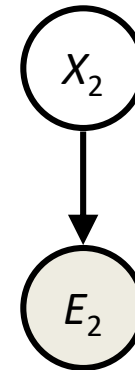
- Every time step, we start with current $P(X \mid \text{evidence})$
- Step 1: Passage of Time (Elapse Time):

$$P(x_t | e_{1:t-1}) = \sum_{x_{t-1}} P(x_{t-1} | e_{1:t-1}) \cdot P(x_t | x_{t-1})$$



- Step 2: Observe:

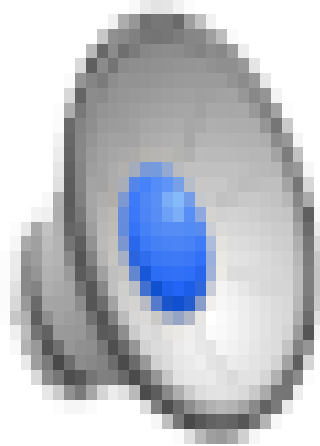
$$P(x_t | e_{1:t}) \propto_X P(x_t | e_{1:t-1}) \cdot P(e_t | x_t)$$



Pacman – Sonar (P4)



Ghost Pacman – Sonar (with beliefs)



Example: Passage of Time

- As time passes, uncertainty “accumulates”

(Transition model: ghosts usually go clockwise)

<0.01	<0.01	<0.01	<0.01	<0.01	<0.01
<0.01	<0.01	<0.01	<0.01	<0.01	<0.01
<0.01	<0.01	1.00	<0.01	<0.01	<0.01
<0.01	<0.01	<0.01	<0.01	<0.01	<0.01

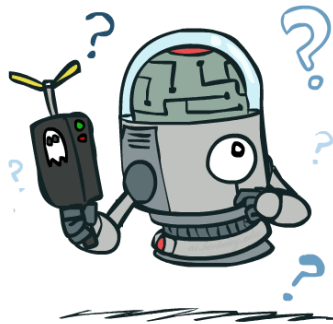
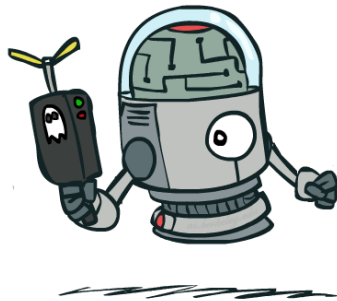
T = 1

<0.01	<0.01	<0.01	<0.01	<0.01	<0.01
<0.01	<0.01	0.06	<0.01	<0.01	<0.01
<0.01	0.76	0.06	0.06	<0.01	<0.01
<0.01	<0.01	0.06	<0.01	<0.01	<0.01

T = 2

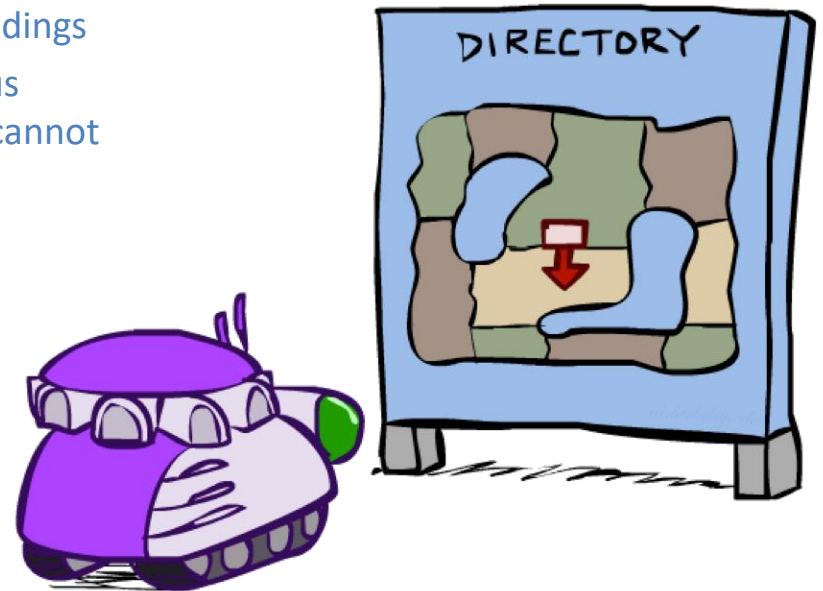
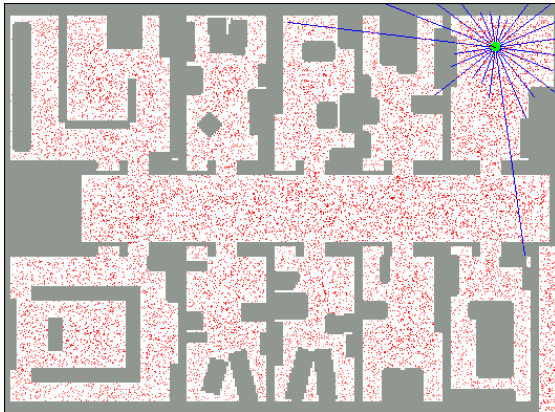
0.05	0.01	0.05	<0.01	<0.01	<0.01
0.02	0.14	0.11	0.35	<0.01	<0.01
0.07	0.03	0.05	<0.01	0.03	<0.01
0.03	0.03	<0.01	<0.01	<0.01	<0.01

T = 5



Robot Localization

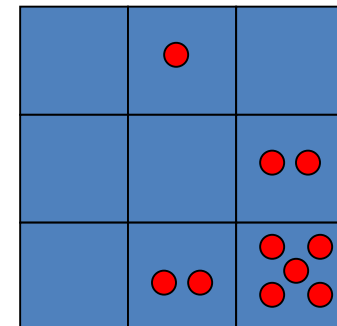
- In robot localization:
 - We know the map, but not the robot's position
 - Observations may be vectors of range finder readings
 - State space and readings are typically continuous (works basically like a very fine grid) and so we cannot store $B(X)$
 - Particle filtering is a main technique



Approximate Inference: Particle Filtering

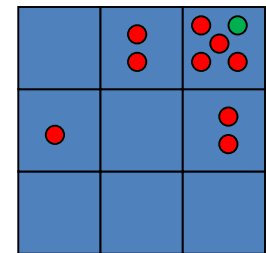
- Filtering: need approximate solution
- Sometimes $|X|$ is too big to use exact inference
 - $|X|$ may be too big to even store $B(X)$
 - E.g. X is continuous
- Solution: approximate inference
 - Track samples of X , not all values
 - Samples are called particles
 - Time per step is linear in the number of samples
 - But: number needed may be large
 - In memory: list of particles, not states
- This is how robot localization works in practice
- Particle is just new name for sample

0.0	0.1	0.0
0.0	0.0	0.2
0.0	0.2	0.5



New Representation: Particles

- Our representation of $P(X)$ is now a list of N particles (samples)
 - Generally, $N \ll |X|$
 - Storing map from X to counts would defeat the point
- $P(x)$ approximated by number of particles with value x
 - So, many x may have $P(x) = 0$!
 - More particles, more accuracy
- For now, all particles have a weight of 1



Particles:

(3,3)
(2,3)
(3,3)
(3,2)
(3,3)
(3,2)
(1,2)
(3,3)
(3,3)
(2,3)

Particle Filtering 1: Elapse Time

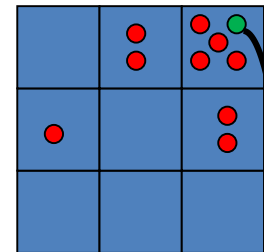
- Each particle is moved by sampling its next position from the transition model

$$x' = \text{sample}(P(X'|x))$$

- This is like prior sampling – samples' frequencies reflect the transition probabilities
 - Here, most samples move clockwise, but some move in another direction or stay in place
- This captures the passage of time**
 - If enough samples, close to exact values before and after (consistent)

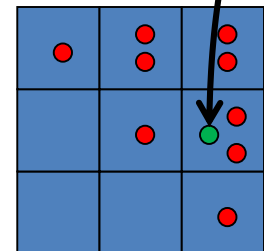
Particles:

(3,3)
(2,3)
(3,3)
(3,2)
(3,3)
(3,2)
(1,2)
(3,3)
(3,3)
(2,3)



Particles:

(3,2)
(2,3)
(3,2)
(3,1)
(3,3)
(3,2)
(1,3)
(2,3)
(3,2)
(2,2)



Particle Filtering 2: Observe

- Slightly trickier:

- Don't sample observation, fix it
- Similar to likelihood weighting, downweight samples based on the evidence

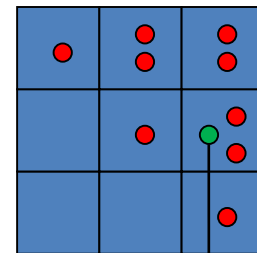
$$w(x) = P(e|x)$$

$$B(X) \propto P(e|X)B'(X)$$

- Note: probabilities don't sum to 1, since all have been downweighted (in fact they now sum to (N times) an approximation of P(e))

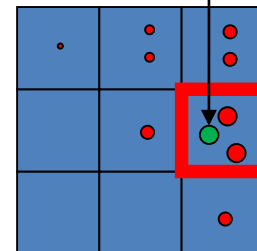
Particles:

(3,2)
(2,3)
(3,2)
(3,1)
(3,3)
(3,2)
(1,3)
(2,3)
(3,2)
(2,2)



Particles:

(3,2) w=.9
(2,3) w=.2
(3,2) w=.9
(3,1) w=.4
(3,3) w=.4
(3,2) w=.9
(1,3) w=.1
(2,3) w=.2
(3,2) w=.9
(2,2) w=.4



Particle Filtering 3: Resample

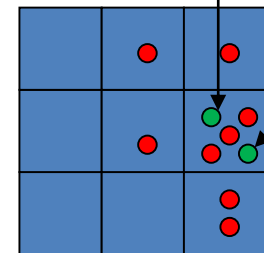
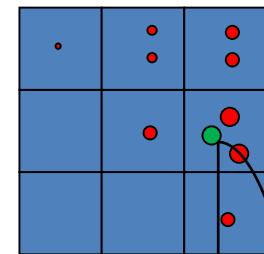
- Instead of tracking weighted samples, we **resample**
- **N times, we choose 1 particle from our weighted sample distribution** (i.e. draw with replacement)
- This is **equivalent to renormalizing** the distribution
- Now the **update is complete** for this time step, continue with the next one

Particles:

(3,2) $w=.9$
(2,3) $w=.2$
(3,2) $w=.9$
(3,1) $w=.4$
(3,3) $w=.4$
(3,2) $w=.9$
(1,3) $w=.1$
(2,3) $w=.2$
(3,2) $w=.9$
(2,2) $w=.4$

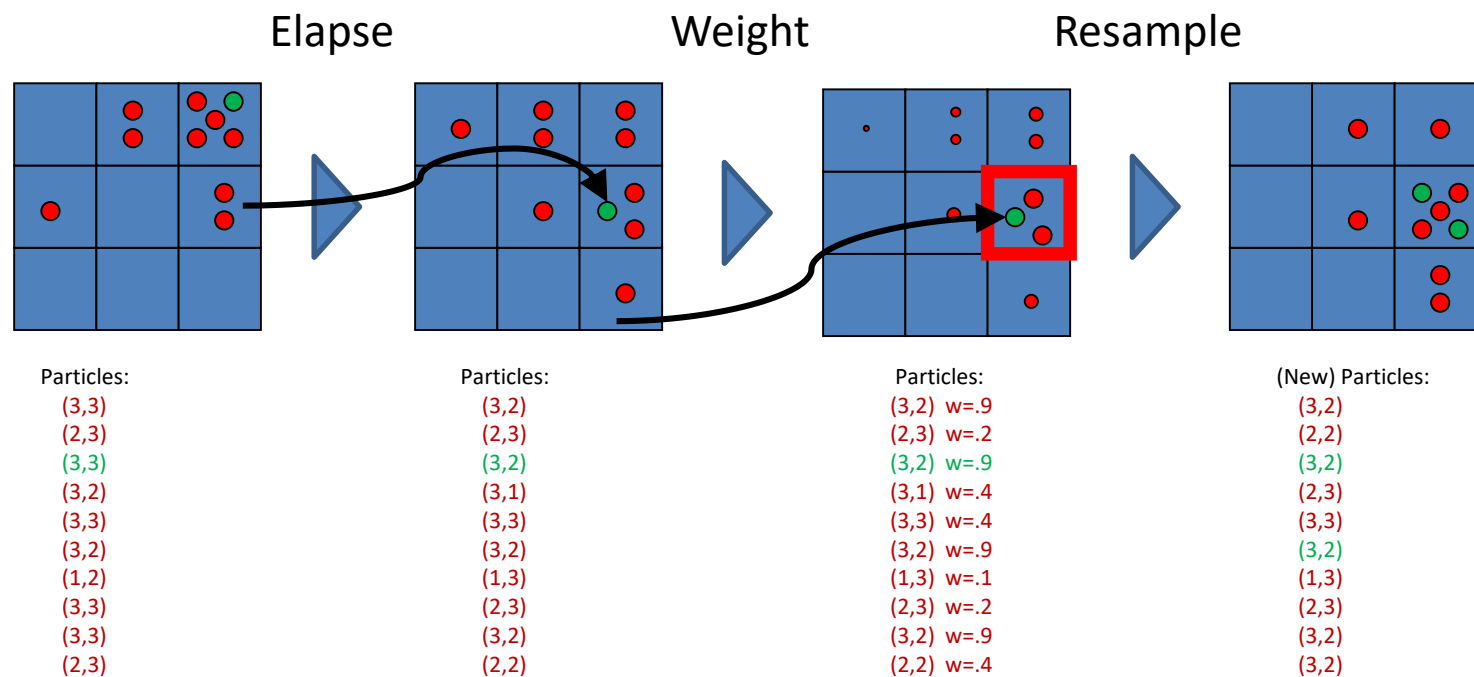
(New) Particles:

(3,2)
(2,2)
(3,2)
(2,3)
(3,3)
(3,2)
(1,3)
(2,3)
(3,2)
(3,2)




Summary: Particle Filtering

- Particles: track samples of states rather than an explicit distribution



Robot Localization (Sonar)



**Global localization with
sonar sensors**

40000

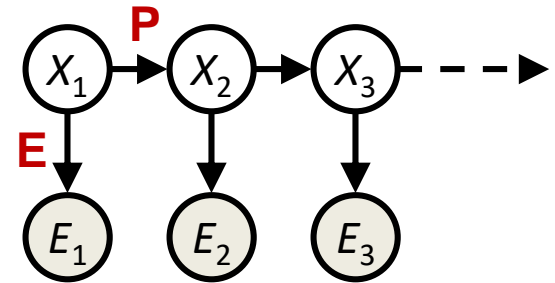
Project 3: Ghostbusters!

- Due **Wed March 21st**
- www.mathcs.emory.edu/~eugene/cs425/p3/
 - **Midterm review (optional): 5:30pm, W301**
 - **Midterm Exam (March 8th)**
 - **Closed book, notes (1 sheet of notes allowed)**

Remaining HMM Issues (after break)

- Where do **P**, **E** probabilities come from?

- Training HMMs
- Dealing with unobserved values



- What's the most likely explanation (a.k.a. decoding)
 - Needed for speech recognition, language parsing, ...
 - Solution: Viterbi algorithm (most likely explanation)
 - What if we want multiple explanations? (beam search)

Techniques so far (on Exam)

- Search
 - Uninformed (UCS)
 - Informed (A^*)
 - Scalability: ID, RM
- Adversarial Search
 - MiniMax
 - Alpha-Beta Pruning
 - ExpectiMax
- Uncertainty:
 - Bayes Nets, Inference
 - Markov Chains, Hidden Markov Models (filtering/prediction)