


# Reinforcement Learning: scaling up

With many slides from Dan Klein and Pieter Abbeel and Stuart Russel

# Review: “Active” Reinforcement Learning

- Want: optimal policy
  - You don’t know the transitions  $T(s,a,s')$
  - You don’t know the rewards  $R(s,a,s')$
  - You choose the actions now
  - Goal: learn the optimal policy / values
- In this case:
  - Learner makes choices!
  - Fundamental tradeoff: **exploration** vs. **exploitation**
  - This is NOT offline planning! **take actions** in the world and get rewards



# Common Confusion

State need not be solely the current sensor readings

- Markov Assumption

Value of state is independent of path taken to reach that state

- Can have memory of the past

Can always create Markovian task by remembering entire past history

# Need for Memory: Simple Example

“out of sight, but not out of mind”

T=1

learning agent



opponent



T=2

learning agent



opponent



Seems reasonable to  
remember opponent  
recently seen

# New Algorithm: Q-Learning

- Q-Learning: **sample-based** Q-value iteration


$$Q_{k+1}(s, a) \leftarrow \sum_{s'} T(s, a, s') \left[ R(s, a, s') + \gamma \max_{a'} Q_k(s', a') \right]$$

- Learn  $Q(s, a)$  values as you go
  - Receive a **sample**  $(s, a, s', r)$
  - Consider your old estimate:  $Q(s, a)$
  - Consider your new sample estimate:

$$sample = R(s, a, s') + \gamma \max_{a'} Q(s', a')$$

- Incorporate new estimate in running average:

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + (\alpha) [sample]$$



# Q-Learning: Implementation Details

Remember, conceptually we are filling in a huge table

		States				
		S0	S1	S2	...	Sn
A c t i o n s	a			.		
	b			.		
	c	...		$Q(S2, c)$		
	.					
	.					
	z					

Tables are a very verbose representation of a function

# Q-Learning: PseudoCode

$$sample = R(s, a, s') + \gamma \max_{a'} Q(s', a')$$

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + (\alpha) [sample]$$

initialize  $Q[S, A]$  arbitrarily

observe current state  $s$

**repeat forever:**

select and carry out an action  $a$

observe reward  $r$  and state  $s'$

$$Q[s, a] \leftarrow Q[s, a] + \alpha (r + \gamma \max_{a'} Q[s', a'] - Q[s, a])$$

$$s \leftarrow s'$$

<https://github.com/aimacode/aima-python/blob/master/rl.py>

# Why does Q-Learning Work?

Jude Shavlik, David Page, Wisconsin

- Intuition: Q-Learning performs iterative approximation
- Each round gets closer to “true” q-value
- If states visited infinitely often, will get infinitely close to true value



# Q-Learning: Convergence Proof

- Applies to Q tables and deterministic, Markovian worlds. Initialize Q's 0 or random finite.
- **Theorem**: if every state-action pair visited infinitely often,  $0 \leq \gamma < 1$ , and  $|\text{rewards}| \leq C$  (some constant), then

$\forall s, a$

$$\lim_{t \rightarrow \infty} \hat{Q}_t(s, a) = Q_{actual}(s, a)$$

the approx. Q table ( $\hat{Q}$ )

the true Q table ( $Q$ )

# Q-Learning Convergence Proof (cont.)

- Consider the max error in the approx. Q-table at step  $t$ :

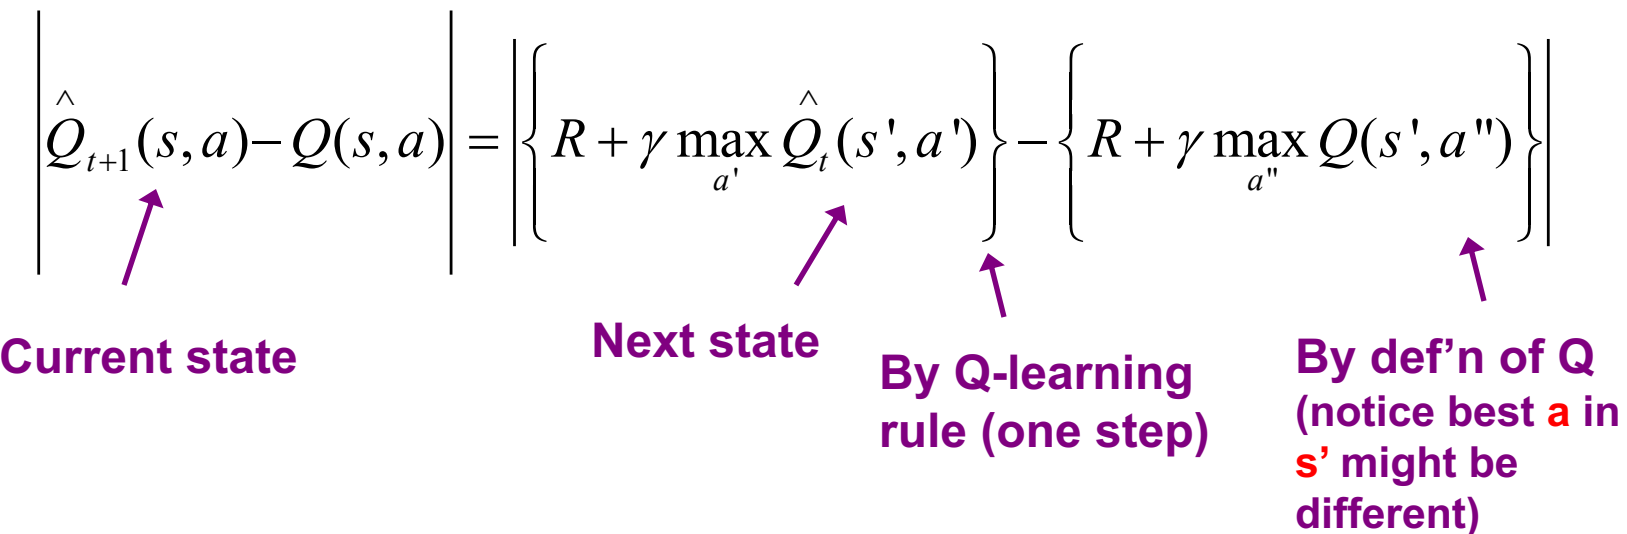
$$\Delta_t \equiv \max_{s,a} | \hat{Q}_t(s,a) - Q_{actual}(s,a) |$$

- The max  $|Q_{actual}(s,a)|$  is finite since  $|r| \leq C$ ,  
so  $\max |Q_{actual}| \leq \sum_{i=0}^{\infty} \gamma^i C = \frac{C}{1-\gamma}$
- Since  $|\hat{Q}_0|$  finite, we have.  $\Delta_0$  **finite**,  
i.e. initial max error is finite

# Q-Learning Convergence Proof (cont.)

Let  $s'$  be the state that results from doing action  $a$  in state  $s$ . Consider what happens when we visit  $s$  and do  $a$  at step  $t + 1$ :

$$\left| \hat{Q}_{t+1}(s, a) - Q(s, a) \right| = \left| \left\{ R + \gamma \max_{a'} \hat{Q}_t(s', a') \right\} - \left\{ R + \gamma \max_{a''} Q(s', a'') \right\} \right|$$



Current state

Next state

By Q-learning rule (one step)

By def'n of Q (notice best  $a$  in  $s'$  might be different)

# Q-Learning Convergence Proof (cont.)

$$= \gamma \left| \max_{a'} \hat{Q}_t(s', a') - \max_{a''} Q(s', a'') \right|$$

By algebra

$$\leq \gamma \max_{a'''} \left| \hat{Q}_t(s', a''') - Q(s', a''') \right|$$

Trickiest step, can  
prove by contradiction

Since  $\left| \max_a f_1(a) - \max_{a'} f_2(a') \right| \leq \max_a \left| f_1(a) - f_2(a) \right|$

$$\leq \gamma \max_{s'', a'''} \left| \hat{Q}_t(s'', a''') - Q(s'', a''') \right|$$

Max at  $s' \leq$  max at any  $s$

$$= \gamma \Delta_t \quad \text{Plugging in defn of } \Delta_t$$

# Q-Learning Convergence Proof (cont.)

- Hence, every time, after  $t$ , we visit an  $\langle s, a \rangle$ , its Q value differs from the correct answer by no more than  $\gamma \Delta_t$
- Let  $T_0 = t_0$  (i.e. the start) and  $T_N$  be the first time since  $T_{N-1}$  where every  $\langle s, a \rangle$  visited at least once
- Call the time between  $T_{N-1}$  and  $T_N$ , a complete interval

Clearly  $\Delta_{T_N} \leq \gamma \Delta_{T_{N-1}}$


# Q-Learning Convergence Proof (concluded)

- That is, every complete interval,  $\Delta_t$  is reduced by at least  $\gamma$
- Since we assumed every  $\langle s, a \rangle$  pair visited infinitely often, we will have an infinite number of complete intervals

$$\text{Hence, } \lim_{t \rightarrow \infty} \Delta_t = 0$$

# Q-Learning Properties

- Amazing result: Q-learning converges to optimal policy -- even if you're acting **suboptimally**!
- This is called **off-policy learning**
- Caveats:
  - You have to explore enough
  - You have to eventually make the learning rate small enough
  - ... but not decrease it too quickly
  - Basically, in the limit, it doesn't matter how you select actions (!)



# Table-Based (Dictionary) Q-Learning

Remember, conceptually we are filling in a huge table



		States				
		S0	S1	S2	...	Sn
Actions	a			.		
	b			.		
	c	...		$Q(S2, c)$		
	.					
	.					
	z					

Tables are a very verbose representation of a function



# Generalizing Across States

- Basic Q-Learning keeps a **table** of all q-values
- In realistic situations, we cannot possibly learn about every single state!
  - Too many states to visit them all in training
  - Too many states to hold the q-tables in memory
- Instead, we want to generalize:
  - Learn about some small number of training states from experience
  - Generalize that experience to new, similar situations
  - This is a fundamental idea in machine learning, and we'll see it over and over again




# RL and Function Approximation

- Exact Q-learning infeasible for many real applications due to curse of dimensionality:  $|S \times A|$  table too big.
- Function Approximation (FA) is a way to “lift the curse:”
  - complexity  $D$  of FA needed to capture regularity in environment may be  $\ll |S|$ .
  - no need to sweep thru entire state space: train on  $N$  “plausible” samples and then **generalize** to similar samples drawn from the same distribution.

# Feature-Based Representations

- Solution: describe a state using a vector of features (properties)
  - Features are functions from states to real numbers (often 0/1) that capture important properties of the state
  - Example features:
    - Distance to closest ghost
    - Distance to closest dot
    - Number of ghosts
    - $1 / (\text{dist to dot})^2$
    - Is Pacman in a tunnel? (0/1)
    - ..... etc.
    - Is it the exact state on this slide?
  - Can also describe a q-state (s, a) with features (e.g. action moves closer to food)



# Linear Value Functions

- Using a feature representation, we can write a q function (or value function) for any state using a few weights:

$$V(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

$$Q(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \dots + w_n f_n(s, a)$$

- Advantage: our experience is summed up in a few powerful numbers
- Disadvantage: states may share features but actually be very different in value!

# Estimating Values 1: Gradient Descent

To find a (local) minimum of a real-valued function  $f(x)$ :

- assign an arbitrary value to  $x$
- repeat

$$x \leftarrow x - \eta \frac{df}{dx}$$

where  $\eta$  is the step size

To find a local minimum of real-valued function  $f(x_1, \dots, x_n)$ :

- assign arbitrary values to  $x_1, \dots, x_n$
- repeat:
  - for each  $x_i$

$$x_i \leftarrow x_i - \eta \frac{\partial f}{\partial x_i}$$

# Estimating Values 1: Linear Regression

- A linear function of variables  $x_1, \dots, x_n$  is of the form

$$f^{\bar{w}}(x_1, \dots, x_n) = w_0 + w_1x_1 + \dots + w_nx_n$$

$\bar{w} = \langle w_0, w_1, \dots, w_n \rangle$  are weights. (Let  $x_0 = 1$ ).

- Given a set  $E$  of examples.

Example  $e$  has input  $x_i = e_i$  for each  $i$  and observed value,  $o_e$ :

$$Error_E(\bar{w}) = \sum_{e \in E} (o_e - f^{\bar{w}}(e_1, \dots, e_n))^2$$

- Minimizing the error using gradient descent, each example should update  $w_i$  using:

$$w_i \leftarrow w_i - \eta \frac{\partial Error_E(\bar{w})}{\partial w_i}$$

# Approximate Q-Learning

$$Q(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \dots + w_n f_n(s, a)$$

- Q-learning with linear Q-functions:

transition =  $(s, a, r, s')$

Error

$$\text{difference} = \left[ r + \gamma \max_{a'} Q(s', a') \right] - Q(s, a)$$


$$Q(s, a) \leftarrow Q(s, a) + \alpha [\text{difference}] \quad \text{Exact Q's}$$

Gradient  
for  $w_i$

$$w_i \leftarrow w_i + \alpha [\text{difference}] f_i(s, a) \quad \text{Approximate Q's}$$

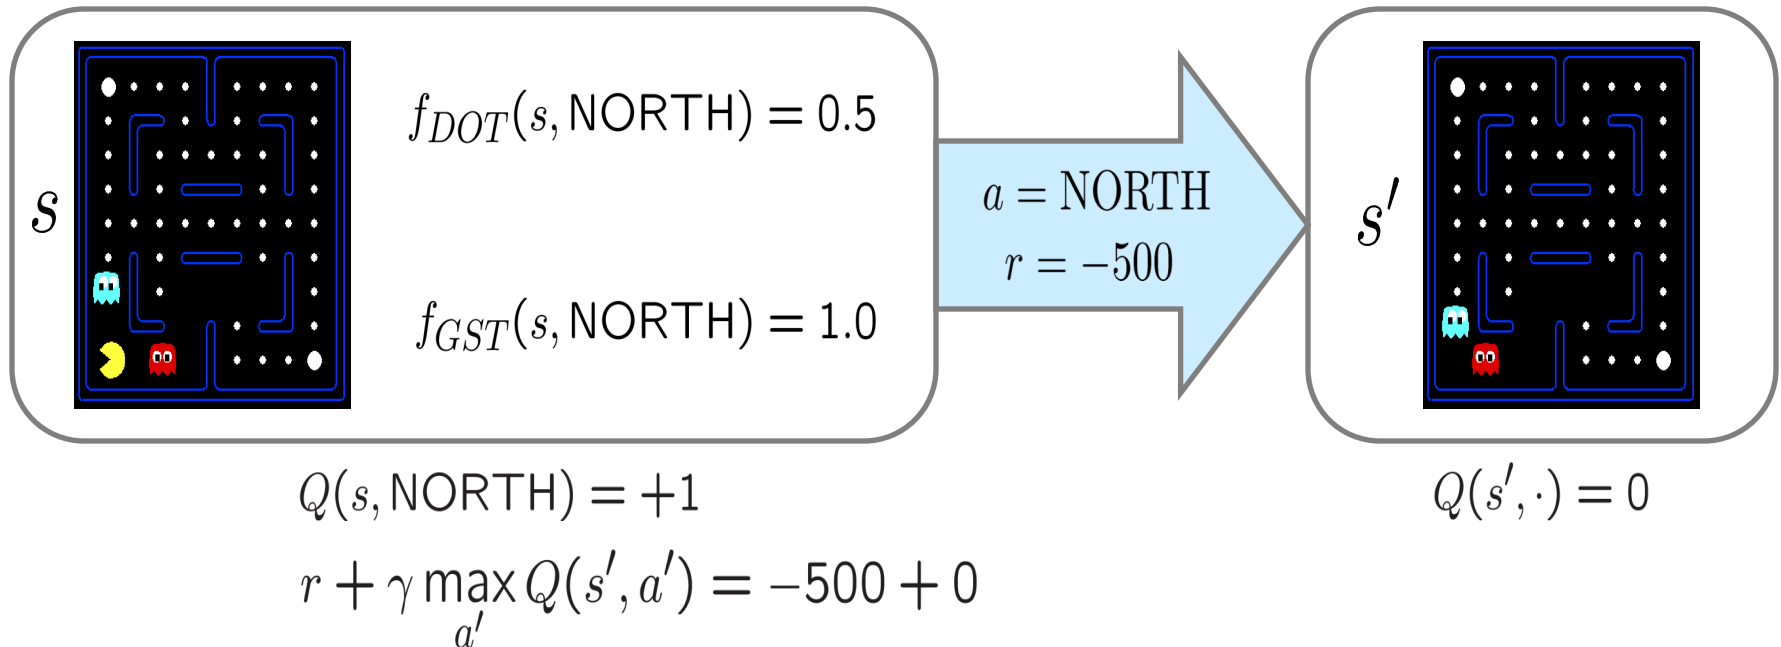
- Intuitive interpretation:

- Adjust weights of active features
- E.g., if something unexpectedly bad happens, *blame the features that were on*: dislike all states with that state's features



# Example: Q-Pacman

$$Q(s, a) = 4.0 f_{DOT}(s, a) - 1.0 f_{GST}(s, a)$$



difference = -501

$$w_{DOT} \leftarrow 4.0 + \alpha [-501] 0.5$$

$$w_{GST} \leftarrow -1.0 + \alpha [-501] 1.0$$

$$Q(s, a) = 3.0 f_{DOT}(s, a) - 3.0 f_{GST}(s, a)$$



# Linear Combination of Features (Proj 4)

- Estimate  $Q(S,a)$  as weighted sum of features (e.g., for Pacman, can use exactly same features as in Proj 2):

$$Q(S,a) = a_1*f_1 + a_2*f_2 + \dots + a_k*f_k$$

$$Q(S,b) = b_1*f_1 + b_2*f_2 + \dots + b_k*f_k$$

- Use linear regression to estimate  $w$ 's:
- For each update of  $Q(S,a)$ :
  - Update  $a_1\dots a_k$  s.t.  $\min(\text{MSE})$